



Lab 3: Hyperparameter Tuning

11210IPT 553000

Deep Learning in Biomedical Optical Imaging

2023/10/02



Outlines

- ▶ Hyperparameters
- ▶ Learning Rate Scheduler
- ▶ Homework 2

Hyperparameters - Model

B. Defining Neural Networks in PyTorch

```
import torch.nn as nn
```

```
# Model definition
```


```
model = nn.Sequential(  
    nn.Flatten(),  
    nn.Linear(256*256*1, 256),  
    nn.ReLU(),  
    nn.Linear(256, 1)  
)
```

```
print(model)
```

```
Sequential(  
  (0): Flatten(start_dim=1, end_dim=-1)  
  (1): Linear(in_features=65536, out_features=256, bias=True)  
  (2): ReLU()  
  (3): Linear(in_features=256, out_features=1, bias=True)  
)
```



Fully Connected Layer (FC): 256



Normal / Pneumonia

Hyperparameters

Number of Units/Neurons



Fully Connected Layer (FC): 32

\hat{y}

Number of Layers


Fully Connected Layer (FC): 256

Fully Connected Layer (FC): 256


Fully Connected Layer (FC): 256

Hyperparameters


Activation Function




sigmoid: $a = \frac{1}{1 + e^{-z}}$



\tanh : $a = \frac{e^x - e^{-x}}{e^x + e^{-x}}$



Relu: $a = \max(0, z)$



Leaky Relu: $a = \max(0.01z, z)$

Hyperparameters

- ▶ Dropout
- ▶ Batch Normalization
- ▶ Regularization (L1, L2)
- ▶

Various layers you can add !

TORCH.NN

These are the basic building blocks for graphs:

`torch.nn`

- Containers
- Convolution Layers
- Pooling layers
- Padding Layers
- Non-linear Activations (weighted sum, nonlinearity)
- Non-linear Activations (other)
- Normalization Layers
- Recurrent Layers
- Transformer Layers
- Linear Layers
- Dropout Layers
- Sparse Layers
- Distance Functions
- Loss Functions
- Vision Layers
- Shuffle Layers
- DataParallel Layers (multi-GPU, distributed)
- Utilities
- Quantized Functions
- Lazy Modules Initialization

[torch.nn](https://pytorch.org/docs/stable/nn.html)

Hyperparameters

A. Data Loading and Preprocessing

```
# Create dataloaders
train_loader = DataLoader(train_dataset, batch_size=32, shuffle=True)
val_loader = DataLoader(val_dataset, batch_size=32, shuffle=False)
```

- Batch size

C. Training the Neural Network

```
import torch.optim as optim
from torch.optim.lr_scheduler import CosineAnnealingLR

train_losses = []
val_losses = []
train_accuracies = []
val_accuracies = []


epochs = 30
best_val_loss = float('inf')

# Criterion and Optimizer
criterion = nn.BCEWithLogitsLoss()
optimizer = optim.Adam(model.parameters(), lr=1e-3)
lr_scheduler = CosineAnnealingLR(optimizer, T_max=len(train_loader)*epochs, eta_min=0)
```

- Number of epochs
- Loss function
- Optimizer
- Learning rate
- Learning rate scheduler

Types of Headache

Migraine




Hypertension



Stress

















Tuning Hyperparameters



Learning Rate Scheduler

- ▶ Dynamically adjusts the learning rate during the training process. It helps to optimize training and sometimes escape local minima.
- ▶ Critical role in training machine learning models effectively for several reasons:
 - ▶ Convergence speed
 - ▶ Training stability
 - ▶ Model Performance





[source](#)

```
from torch.optim.lr_scheduler import CosineAnnealingLR
```

```
lr_scheduler = CosineAnnealingLR(optimizer, T_max=len(train_loader)*epochs, eta_min=0)
```

Docs > torch.optim



COSINEANNEALINGLR

`lr_scheduler.PolynomialLR`

Decays the learning rate of each parameter group using a polynomial function in the given total_iters.

`lr_scheduler.CosineAnnealingLR`

Set the learning rate of each parameter group using a cosine annealing schedule, where η_{max} is set to the initial lr and T_{cur} is the number of epochs since the last restart in SGDR:

`lr_scheduler.ChainedScheduler`

Chains list of learning rate schedulers.

`lr_scheduler.SequentialLR`

Receives the list of schedulers that is expected to be called sequentially during optimization process and milestone points that provides exact intervals to reflect which scheduler is supposed to be called at a given epoch.

`lr_scheduler.ReduceLROnPlateau`

Reduce learning rate when a metric has stopped improving.

`lr_scheduler.CyclicLR`

Sets the learning rate of each parameter group according to cyclical learning rate policy (CLR).

`lr_scheduler.OneCycleLR`

Sets the learning rate of each parameter group according to the 1cycle learning rate policy.

CLASS `torch.optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max, eta_min=0, last_epoch=-1, verbose=False)` [SOURCE]

Set the learning rate of each parameter group using a cosine annealing schedule, where η_{max} is set to the initial lr and T_{cur} is the number of epochs since the last restart in SGDR:

$$\eta_t = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min}) \left(1 + \cos \left(\frac{T_{cur}}{T_{max}} \pi \right) \right), \quad T_{cur} \neq (2k+1)T_{max};$$
$$\eta_{t+1} = \eta_t + \frac{1}{2}(\eta_{max} - \eta_{min}) \left(1 - \cos \left(\frac{1}{T_{max}} \pi \right) \right), \quad T_{cur} = (2k+1)T_{max}.$$

When `last_epoch=-1`, sets initial lr as lr. Notice that because the schedule is defined recursively, the learning rate can be simultaneously modified outside this scheduler by other operators. If the learning rate is set solely by this scheduler, the learning rate at each step becomes:

$$\eta_t = \eta_{min} + \frac{1}{2}(\eta_{max} - \eta_{min}) \left(1 + \cos \left(\frac{T_{cur}}{T_{max}} \pi \right) \right)$$

It has been proposed in **SGDR: Stochastic Gradient Descent with Warm Restarts**. Note that this only implements the cosine annealing part of SGDR, and not the restarts.



Parameters:

- **optimizer** (*Optimizer*) – Wrapped optimizer.
- **T_max** (*int*) – Maximum number of iterations.
- **eta_min** (*float*) – Minimum learning rate. Default: 0.
- **last_epoch** (*int*) – The index of last epoch. Default: -1.
- **verbose** (*bool*) – If `True`, prints a message to stdout for each update. Default: `False`.

<https://pytorch.org/docs/stable/optim.html>



CosineAnnealingLR

CosineAnnealingLR





```
model = nn.Sequential(  
    nn.Flatten(),  
    nn.Linear(256*256*1, 256),  
    nn.ReLU(),  
    nn.Linear(256, 1)  
).cuda()
```

StepLR





CosineAnnealingLR


```
model = nn.Sequential(  
    nn.Flatten(),  
  
    nn.Linear(256*256*1, 64),  
    nn.BatchNorm1d(64),  
    nn.ReLU(),  
    nn.Dropout(0.5),  
  
    nn.Linear(64, 64),  
    nn.BatchNorm1d(64),  
    nn.ReLU(),  
    nn.Dropout(0.5),  
  
    nn.Linear(64, 64),  
    nn.BatchNorm1d(64),  
    nn.ReLU(),  
    nn.Dropout(0.5),  
  
    nn.Linear(64, 1)  
)  
print(model)
```



StepLR



Homework 2

- ▶ **Deadline:** 23:59, 16th Oct. (GMT+8)
- ▶ We have 2 parts, the code and the report. Details are in hw2_description.pdf
- ▶  **Important:** Make sure your commit is timestamped before the deadline. Late submissions might not be graded or could incur a penalty. Only the GitHub link is required on NTHU EEclass.

CODING TIME!!

