



Lab 4: Build a CNN

11210IPT 553000

Deep Learning in Biomedical Optical Imaging


2023/10/16



Outlines

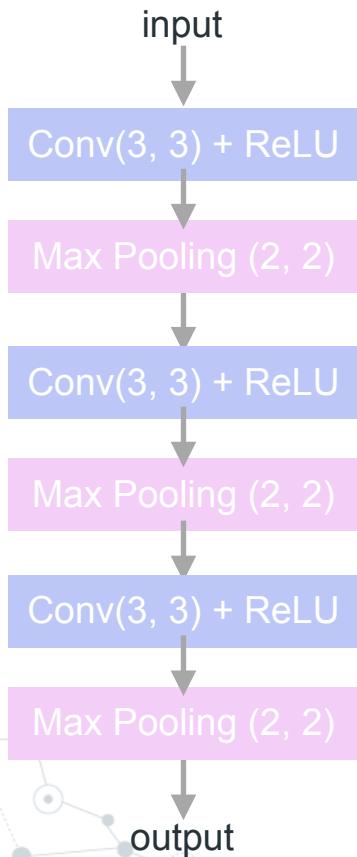
- ▶ Build a Model
- ▶ PyTorch Layers for CNN
- ▶ Global Average Pooling
- ▶ Homework 3

Build a Model - ANN



```
class LinearModel(nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.flatten = nn.Flatten()  
        self.fc1 = nn.Linear(256*256*1, 32)  
        self.fc2 = nn.Linear(32, 32)  
        self.fc3 = nn.Linear(32, 32)  
        self.fc4 = nn.Linear(32, 1)  
  
    def forward(self, x):  
        x = self.flatten(x)  
        x = F.relu(self.fc1(x))  
        x = F.relu(self.fc2(x))  
        x = F.relu(self.fc3(x))  
        return self.fc4(x)
```

Build a Model - CNN



```
class ConvModel(nn.Module):
    def __init__(self):
        super().__init__()

        # 1 channel, and using 3x3 kernels for simplicity, 256*256
        self.conv1 = nn.Conv2d(1, 32, kernel_size=3, stride=1, padding='same')
        self.pool1 = nn.MaxPool2d(kernel_size=2, stride=2) # 128*128

        self.conv2 = nn.Conv2d(32, 32, kernel_size=3, stride=1, padding='same') # 128*128
        self.pool2 = nn.MaxPool2d(kernel_size=2, stride=2) # 64*64

        self.conv3 = nn.Conv2d(32, 32, kernel_size=3, stride=1, padding='same') # 64*64
        self.pool3 = nn.MaxPool2d(kernel_size=2, stride=2) # 32*32

        # Adjust flattened dimensions based on the output size of your last pooling layer
        flattened_dim = 32 * 32 * 32

        self.fc1 = nn.Linear(flattened_dim, 32)
        self.fc2 = nn.Linear(32, 1)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = self.pool1(x)

        x = F.relu(self.conv2(x))
        x = self.pool2(x)

        x = F.relu(self.conv3(x))
        x = self.pool3(x)

        # Flatten the output for the fully connected layers
        x = x.reshape(x.size(0), -1) # x.size(0) is the batch size

        x = F.relu(self.fc1(x))
        return self.fc2(x)
```

PyTorch Layers for CNN

Docs > torch.nn



Convolution Layers

`nn.Conv1d`

Applies a 1D convolution over an input signal composed of several input planes.

`nn.Conv2d`

Applies a 2D convolution over an input signal composed of several input planes.

`nn.Conv3d`

Applies a 3D convolution over an input signal composed of several input planes.

`nn.ConvTranspose1d`


Applies a 1D transposed convolution operator over an input image composed of several input planes.

`nn.ConvTranspose2d`

Applies a 2D transposed convolution operator over an input image composed of several input planes.

`nn.ConvTranspose3d`

Applies a 3D transposed convolution operator over an input image composed of several input planes.



<https://pytorch.org/docs/stable/nn.html#convolution-layers>

PyTorch Layers for CNN

Docs > torch.nn



Pooling layers

`nn.MaxPool1d`

Applies a 1D max pooling over an input signal composed of several input planes.

`nn.MaxPool2d`

Applies a 2D max pooling over an input signal composed of several input planes.

`nn.MaxPool3d`

Applies a 3D max pooling over an input signal composed of several input planes.

`nn.MaxUnpool1d`

Computes a partial inverse of `MaxPool1d`.

`nn.MaxUnpool2d`


Computes a partial inverse of `MaxPool2d`.

`nn.MaxUnpool3d`

Computes a partial inverse of `MaxPool3d`.

`nn.AvgPool1d`

Applies a 1D average pooling over an input signal composed of several input planes.



<https://pytorch.org/docs/stable/nn.html#pooling-layers>

Global Average Pooling

Docs > torch.nn > AdaptiveAvgPool2d



ADAPTIVEAVGPOOL2D

CLASS `torch.nn.AdaptiveAvgPool2d(output_size)` [SOURCE]

Applies a 2D adaptive average pooling over an input signal composed of several input planes.

The output is of size $H \times W$, for any input size. The number of output features is equal to the number of input planes.

Parameters

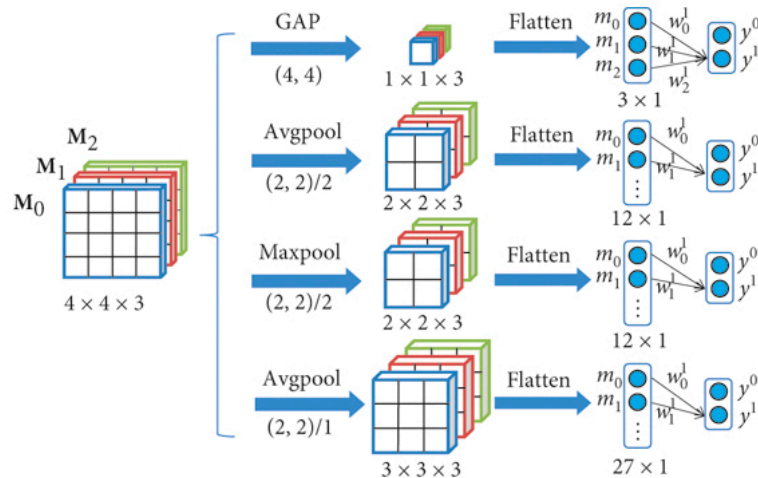
output_size (`Union[int, None, Tuple[Optional[int], Optional[int]]]`) – the target output size of the image of the form $H \times W$. Can be a tuple (H, W) or a single H for a square image $H \times H$. H and W can be either a `int`, or `None` which means the size will be the same as that of the input.

Shape:

- Input: (N, C, H_{in}, W_{in}) or (C, H_{in}, W_{in}) .
- Output: (N, C, S_0, S_1) or (C, S_0, S_1) , where $S = \text{output_size}$.


Examples

```
>>> # target output size of 5x7
>>> m = nn.AdaptiveAvgPool2d((5, 7))
>>> input = torch.randn(1, 64, 8, 9)
>>> output = m(input)
>>> # target output size of 7x7 (square)
>>> m = nn.AdaptiveAvgPool2d(7)
>>> input = torch.randn(1, 64, 10, 9)
```



<https://pytorch.org/docs/stable/generated/torch.nn.AdaptiveAvgPool2d.html#torch.nn.AdaptiveAvgPool2d>

Homework 3

- ▶ **Deadline:** 23:59, 30th Oct. (GMT+8)
- ▶ We need to write a report to answer questions. Details are in hw3_description.pdf
- ▶  **Important:** Make sure your commit is timestamped before the deadline. Late submissions might not be graded or could incur a penalty. Only the GitHub link is required on NTHU EEclass.

CODING TIME!!

