

11210IPT553000

Deep Learning in Biomedical Optical Imaging

Week 7

Convolutional Neural Network

Foundations of Convolutional Neural Networks

Instructor: Hung-Wen Chen
2023/10/23 @NTHU, Fall 2023

coursera



deeplearning.ai

SPONSORED BY

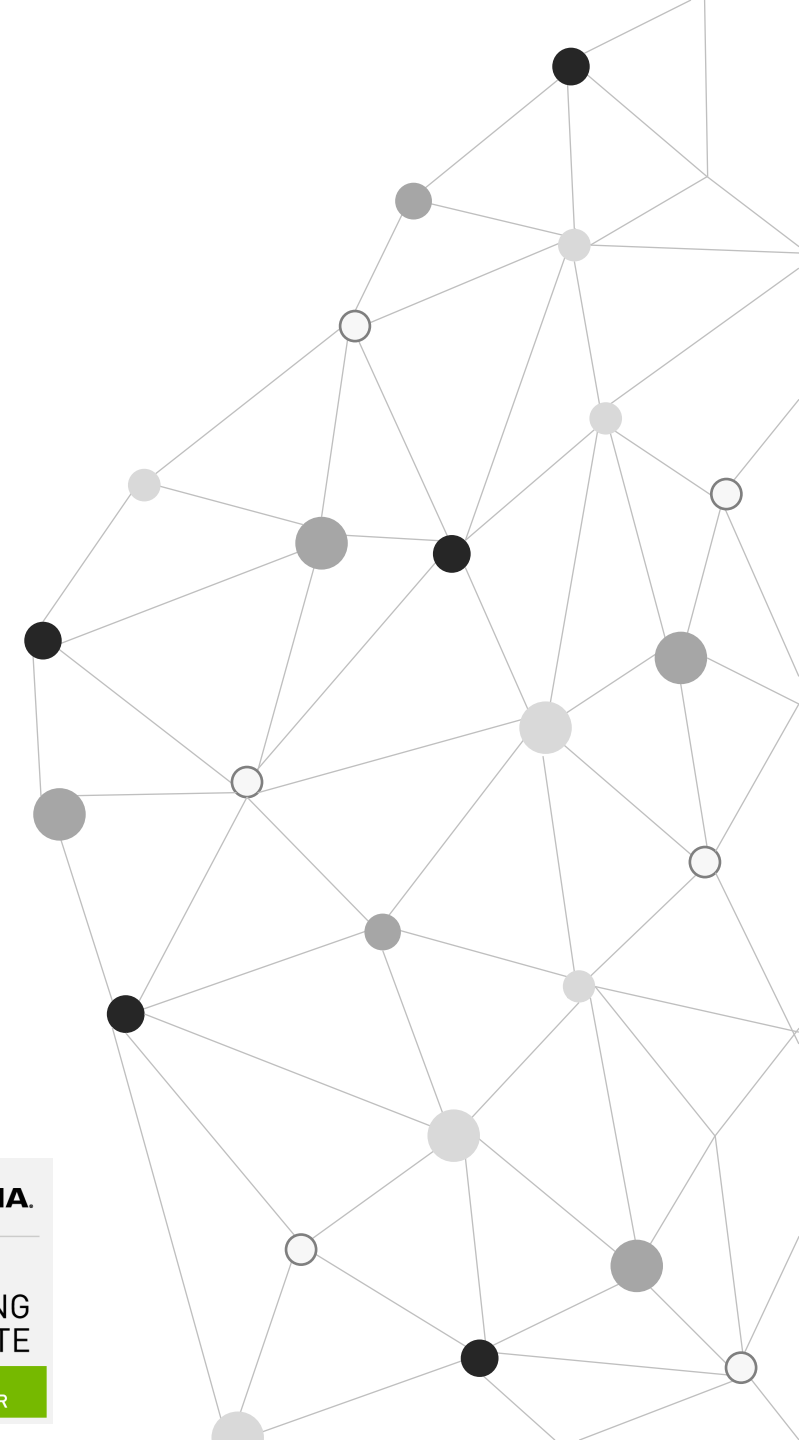


DEEP
LEARNING
INSTITUTE




DEEP
LEARNING
INSTITUTE

CERTIFIED
INSTRUCTOR



- Computer Vision
- Foundations of Convolutional Neural Networks (Course 4 Week 1)
- Lab Practice: Transfer Learning

Computer Vision




Supervised Learning

Structured Data


Age	Weight	...	Gender
50	72		Male
33	82		Female
18	66		Female
⋮	⋮		⋮
80	55		Male

User Age	Ad Id	...	Click
41	93242		1
80	93287		0
18	87312		1
⋮	⋮		⋮
27	71244		1

Unstructured Data



Audio



Four scores and seven
years ago...

Text


Computer Vision

Tasks

Classification



Classification + Localization



CAT

Object Detection




DOG, DOG, CAT

Instance Segmentation



DOG, DOG, CAT

Semantic Segmentation



GRASS, CAT,
TREE, SKY

Single Object


Multiple Object

This image is CC0 public domain


No objects, just pixels

Computer Vision


Instance segmentation vs Semantic segmentation



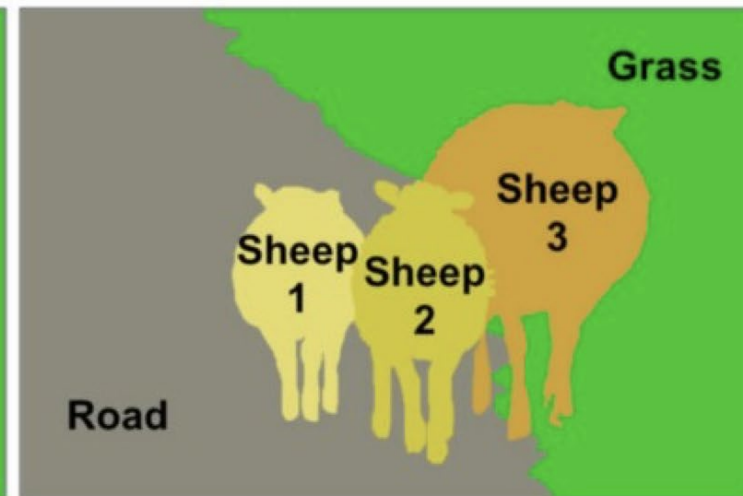
Classification + Localization



Object Detection




Semantic Segmentation



Instance Segmentation

Computer Vision

Deep Learning on large images



64x64 x3


→ Cat? (0/1)

12288




1000 x 1000 x 3
= 3 million

3 billion!!!




Computer Vision

Computer vision problem




vertical edges



horizontal edges

Image Kernels (Filters)



Computer Vision

The representation of an image

1 channel: Black -> 0, White -> 255 // 2 dimensions



12X16

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

157	153	174	168	150	152	129	151	172	161	155	156
155	182	163	74	75	62	33	17	110	210	180	154
180	180	50	14	34	6	10	33	48	106	159	181
206	109	5	124	131	111	120	204	166	15	56	180
194	68	137	251	237	239	239	228	227	87	71	201
172	105	207	233	233	214	220	239	228	98	74	206
188	88	179	209	185	215	211	158	139	75	20	169
189	97	165	84	10	168	134	11	31	62	22	148
199	168	191	193	158	227	178	143	182	106	36	190
205	174	155	252	236	231	149	178	228	43	95	234
190	216	116	149	236	187	86	150	79	38	218	241
190	224	147	108	227	210	127	102	36	101	255	224
190	214	173	66	103	143	96	50	2	109	249	215
187	196	235	75	1	81	47	0	6	217	255	211
183	202	237	145	0	0	12	108	200	138	243	236
195	206	123	207	177	121	123	200	175	13	96	218

Computer Vision

Vertical edge detection

$$3 \times 1 + 1 \times 1 + 1 \times 2 + 0 \times 0 + 0 \times 5 + 0 \times 7 + 1 \times -1 + 8 \times -1 + 2 \times -1$$

3 ¹	0 ¹	1 ⁻¹	2 ⁻¹	7 ⁻⁰	4 ⁻¹
1 ¹	5 ¹	8 ⁻¹	9 ⁻¹	3 ⁻⁰	1 ⁻¹
2 ¹	7 ¹	2 ⁻¹	5 ⁻¹	1 ⁻⁰	3 ⁻¹
0 ¹	1 ¹	3 ⁻¹	1 ⁻¹	7 ⁻⁰	8 ⁻¹
4	2	1	6	2	8
2	4	5	2	3	9

6 X 6

Filter or Kernel

1	0	-1
1	0	-1
1	0	-1

3 X 3

*

=

-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

4 X 4

Convolution

Computer Vision

Kernels and Convolution

Blur Kernel

.06	.13	.06
.13	.25	.13
.06	.13	.06


*

Original Image

1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

=

Convolved Image



Computer Vision

Kernels and Convolution

Blur Kernel

.06	.13	.06
.13	.25	.13
.06	.13	.06

*

Original Image

1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

=

Convolved Image

Computer Vision

Kernels and Convolution

Blur Kernel

Original Image

Convolved Image

Multiply

.06	.13	.06
.13	.25	.13
.06	.13	.06

*

.06	0	.06	1	0	1
0	.25	0	0	1	0
0	.13	.06	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

=

Computer Vision

Kernels and Convolution

Blur Kernel

.06	.13	.06
.13	.25	.13
.06	.13	.06

*

Original Image

.06	0	.06	1	0	1
0	.25	0	0	1	0
0	.13	.06	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

Total

=

Convolved Image

.56

Computer Vision

Kernels and Convolution

Blur Kernel

.06	.13	.06
.13	.25	.13
.06	.13	.06

*

Original Image

1	0	.13	.06	0	1
0	.13	0	0	1	0
0	.06	.13	.06	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

=

Convolved Image

.56	.57		

Computer Vision

Kernels and Convolution

Blur Kernel

.06	.13	.06
.13	.25	.13
.06	.13	.06

*

Original Image

1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

=

Convolved Image

.56	.57	.57	.56
.7	.82	.82	.7
.69	.95	.95	.69
.64	.69	.69	.64

1 _{x1}	1 _{x0}	1 _{x1}	0	0
0 _{x0}	1 _{x1}	1 _{x0}	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved
Feature

1	0	1
0	1	0
1	0	1

Filter or Kernel

Computer Vision

Vertical edge detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

6 X 6

*

1	0	-1
1	0	-1
1	0	-1

3 X 3



=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

4 X 4




*



Computer Vision


Vertical edge detection examples

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0




*

1	0	-1
1	0	-1
1	0	-1




=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0




0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10




*

1	0	-1
1	0	-1
1	0	-1



=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0



Vertical and horizontal edge detection

1	0	-1
1	0	-1
1	0	-1

Vertical

1	1	1
0	0	0
-1	-1	-1

Horizontal


10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

*

1	1	1
0	0	0
-1	-1	-1

=

0	0	0	0
30	10	-10	-30
30	10	-10	-30
0	0	0	0



Computer Vision

Learning to detect edges

1	0	-1
1	0	-1
1	0	-1

1	0	-1
2	0	-2
1	0	-1

3	0	-3
10	0	-10
3	0	-3


Sobel Filter

Scharr Filter

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

*

w_1	w_2	w_3
w_4	w_5	w_6
w_7	w_8	w_9




=

45 degree
60 degree
90 degree
.....


Computer Vision

Edge detection example


Original Image



Vertical Edges



Horizontal Edges



0	0	0
0	1	0
0	0	0

1	0	-1
2	0	-2
1	0	-1

1	2	1
0	0	0
-1	-2	-1

Computer Vision

Kernels and Convolution



Original Image



Blur



Sharpen



Brighten



Darken



Computer Vision

Kernels and Convolution




Original Image





Blur

.06	.13	.06
.13	.25	.13
.06	.13	.06




Sharpen

0	-1	0
-1	5	-1
0	-1	0




Brighten

0	0	0
0	1.5	0
0	0	0



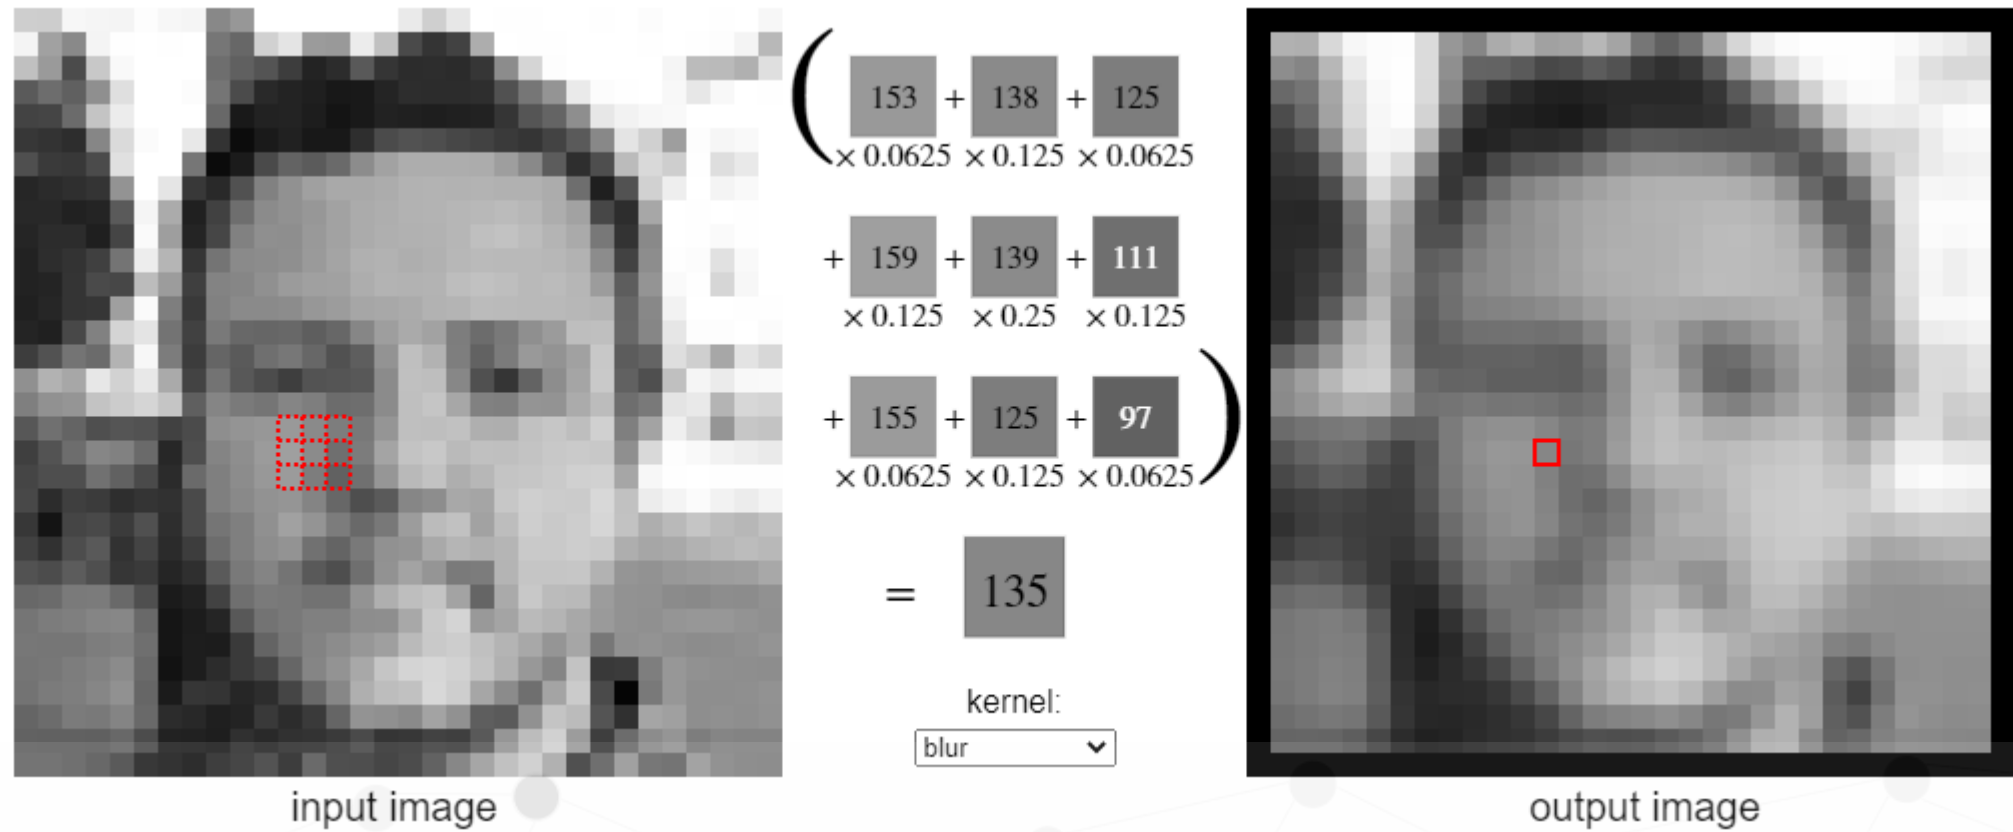
Darken

0	0	0
0	0.5	0
0	0	0



Computer Vision

Image kernels (filters) explained visually



Computer Vision

Image kernels (filters) explained visually

0	0	0
0	1	0
0	0	0

identity ▼

-1	-1	-1
-1	8	-1
-1	-1	-1

outline ▼

-2	-1	0
-1	1	1
0	1	2

emboss ▼




<https://setosa.io/ev/image-kernels/>

Convolution is more often taken over an infinite range,

$$f * g \equiv \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$$

$$= \int_{-\infty}^{\infty} g(\tau) f(t - \tau) d\tau$$



The animations above graphically illustrate the convolution of two [boxcar functions](#) (left) and two [Gaussians](#) (right). In the plots, the green curve shows the convolution of the blue and red curves as a function of t , the position indicated by the vertical green line. The gray region indicates the product $g(\tau) f(t - \tau)$ as a function of t , so its area as a function of t is precisely the convolution. One feature to emphasize and which is not conveyed by these illustrations (since they both exclusively involve symmetric functions) is that the function g must be mirrored before lagging it across f and integrating. (<https://mathworld.wolfram.com/Convolution.html>)

The name Convolutional Neural Networks (CNN) suggests that they use the convolution operation, but in the usual way to describe CNN, **it is correlation that it's using**. However, convolution and correlation can be interchanged through a simple rotation operation. Therefore, the name Convolutional Neural Networks is also justified.

Technical Note on Cross-Correlation vs. Convolution

Convolution in math textbook

2	3	7	4	6	2
6	6	9	8	7	4
3	4	8	3	8	9
7	8	3	6	6	3
4	2	1	8	3	4
3	2	4	1	9	8


*

3	4	5
1	0	2
-1	9	7

7	9	-1
2	0	1
5	4	3


Mirroring it both on the vertical and horizontal axes

Padding




Padding

Output dimension without padding




6 X 6

*



3 X 3

=



4 X 4

Generalize:


$n \times n$

$f \times f$

$(n-f+1) \times (n-f+1)$


Padding

Output dimension with padding




6 X 6

*



3 X 3

=



4 X 4

6 X 6

$p = 1$

8 X 8

Generalize: $(n+2p) \times (n+2p)$

$f \times f$

$(n+2p-f+1) \times (n+2p-f+1)$

No padding

“Valid”:

$$(n \times n) * (f \times f) = (n - f + 1) \times (n - f + 1)$$
$$(6 \times 6) * (3 \times 3) = 4 \times 4$$

“Same”: Pad so that output size is the same as the input size.

$$(n + 2p - f + 1) \times (n + 2p - f + 1) \qquad n + 2p - f + 1 = n \rightarrow p = (f - 1) / 2$$
$$(3 \times 3) \rightarrow p = (3 - 1) / 2 = 1 \qquad (5 \times 5) \rightarrow p = (5 - 1) / 2 = 2$$

The value of f is usually “odd”

Computer Vision

Zero Padding

Original Image

1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

Zero Padding

0	0	0	0	0	0	0	0
0	1	0	1	1	0	1	0
0	0	1	0	0	1	0	0
0	0	1	1	1	1	0	0
0	0	1	1	1	1	0	0
0	1	0	1	1	0	1	0
0	1	1	0	0	1	1	0
0	0	0	0	0	0	0	0

Computer Vision

Mirror Padding


Original Image

1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0
0	1	1	1	1	0
1	0	1	1	0	1
1	1	0	0	1	1

Mirror Padding

1	1	0	1	1	0	1	1
1	1	0	1	1	0	1	1
0	0	1	0	0	1	0	0
0	0	1	1	1	1	0	0
0	0	1	1	1	1	0	0
1	1	0	1	1	0	1	1
1	1	1	0	0	1	1	1
1	1	1	0	0	1	1	1

Strided Convolutions




Strided convolutions

Different strides

Stride 1


1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0



.56	.57	.57	.56
-----	-----	-----	-----

Stride 2


1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0



.56	.57
-----	-----

Stride 3

1	0	1	1	0	1
0	1	0	0	1	0
0	1	1	1	1	0



.56	.56
-----	-----

Strided Convolutions

Example

2 ³	3 ⁴	7 ³	4 ⁴	6 ³	2 ⁴	9 ⁴
6 ¹	6 ⁰	9 ¹	8 ⁰	7 ¹	4 ⁰	3 ²
3 ⁻³	4 ⁴	8 ³	3 ⁴	8 ³	9 ⁴	7 ⁴
7 ¹	8 ⁰	3 ¹	6 ⁰	6 ¹	3 ⁰	4 ²
4 ⁻³	2 ⁴	1 ³	8 ⁴	3 ³	4 ⁴	6 ⁴
3 ¹	2 ⁰	4 ¹	1 ⁰	9 ¹	8 ⁰	3 ²
0 ⁻¹	1 ⁰	3 ⁻¹	9 ⁰	2 ⁻¹	1 ⁰	4 ³

7 X 7

*

3	4	4
1	0	2
-1	0	3

3 X 3

=

3 X 3

Stride s = 2

$\lfloor x \rfloor = \text{floor}(x)$

$n \times n$ image * $f \times f$ filter

padding p

stride s

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$

$$\left\lfloor \frac{7+2*0-3}{2} + 1 \right\rfloor = 3$$

Strided Convolutions

Output dimension


$n \times n$ image $f \times f$ filter padding p stride s

Output size

$$\left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor \times \left\lfloor \frac{n+2p-f}{s} + 1 \right\rfloor$$


$$\lfloor x \rfloor = \text{floor}(x)$$

Convolutions Over Volumes




Convolution Over Volumes

Convolutions on RGB images




$6 \times 6 \times 3$

*



$3 \times 3 \times 3$


=



4×4


Convolution Over Volumes

Convolutions on RGB image




6 X 6 X 3

*




3 X 3 X 3

27 numbers



=



4 X 4

1	0	-1
1	0	-1
1	0	-1

0	0	0
0	0	0
0	0	0

0	0	0
0	0	0
0	0	0


3 X 3 X 3

1	0	-1
1	0	-1
1	0	-1


1	0	-1
1	0	-1
1	0	-1

1	0	-1
1	0	-1
1	0	-1


Convolution Over Volumes




Multiple filters
Vertical edge




Horizontal edge





=



=



4 x 4 x 2



Summary: $n \times n \times n_c * f \times f \times n_c \rightarrow (n-f+1) \times (n-f+1) \times n'_c$ # of filters

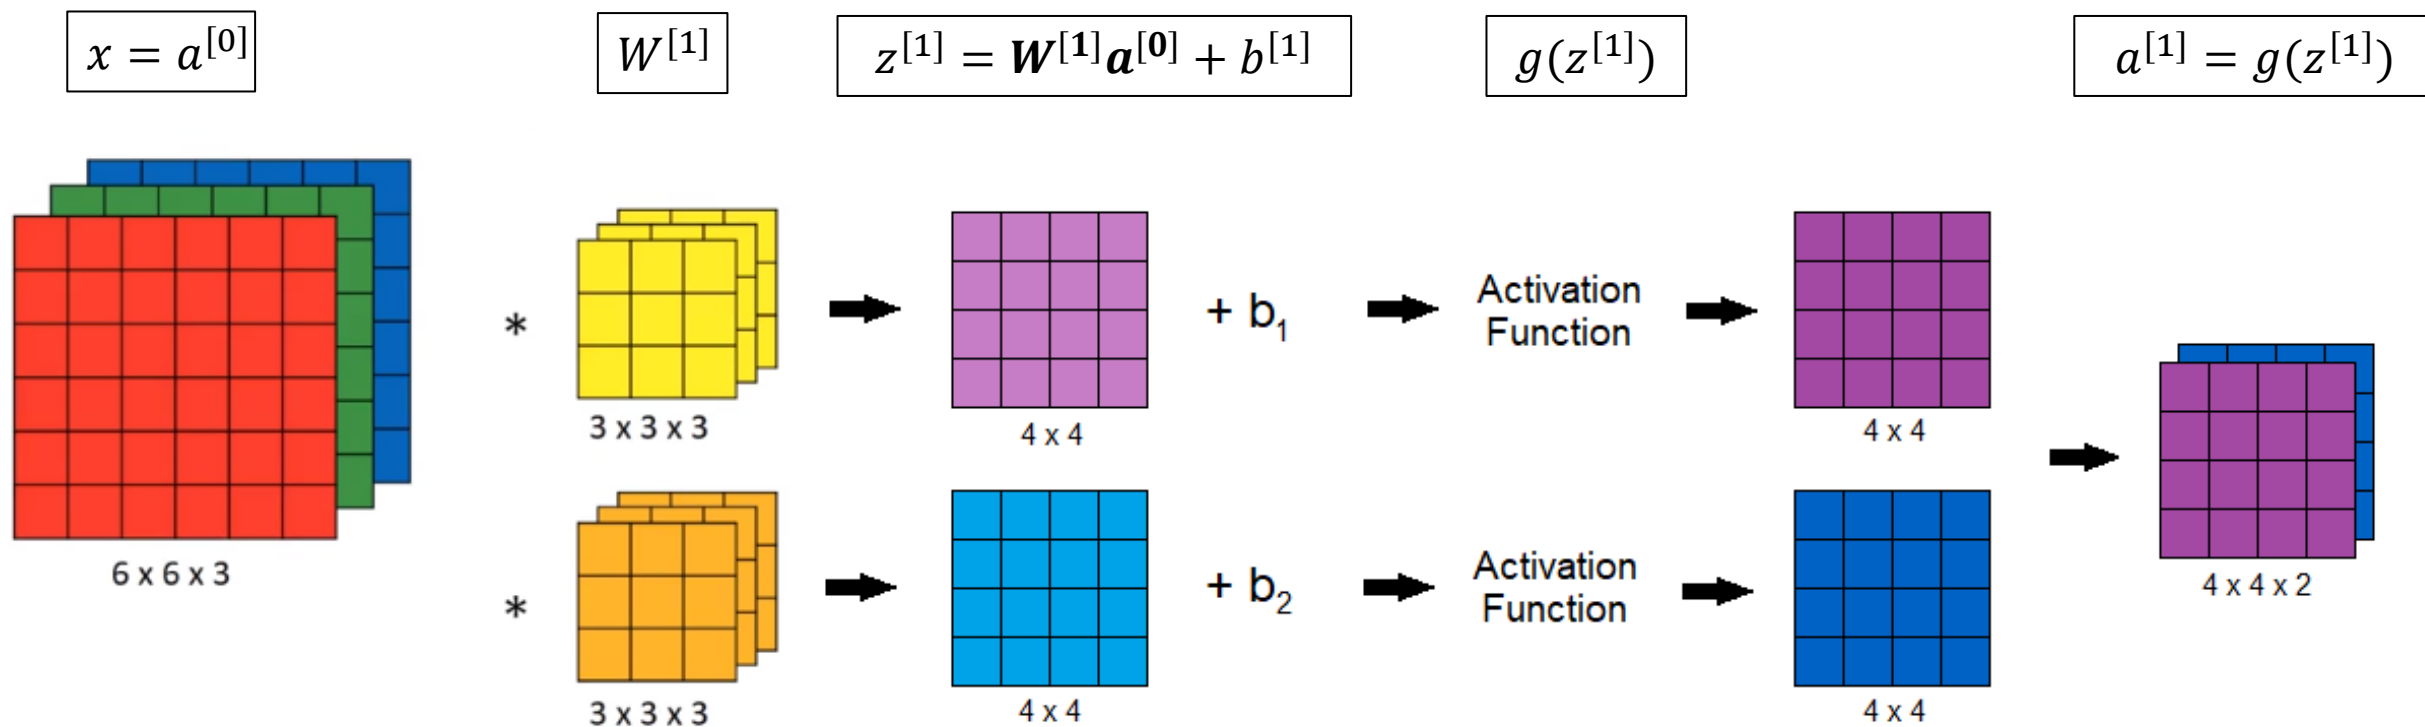
One Layer of a Convolutional Network



Yann LeCun

One Layer of a Convolutional Network


Example of a layer



One Layer of a Convolutional Network

Number of parameters in one layer

If you have 10 filters that $3 \times 3 \times 3$ in one layer of a neural network, how many parameters does this layer have?



27 parameters + 1 bias
= 28 parameters

$28 \times 10 = 280$ parameters

If layer l is a convolutional layer:

$f^{[l]}$: filter size in layer l

$p^{[l]}$: padding in layer l

$s^{[l]}$: stride in layer l

$n_c^{[l]}$: # of filters

Each filter is: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]}$

Activation $a^{[l]} : n_H^{[l]} \times n_w^{[l]} \times n_c^{[l]}$

Weights: $f^{[l]} \times f^{[l]} \times n_c^{[l-1]} \times n_c^{[l]}$

Bias: $n_c^{[l]}$ or $(1,1,1, n_c^{[l]})$


Input: $n_H^{[l-1]} \times n_w^{[l-1]} \times n_c^{[l-1]}$

Output: $n_H^{[l]} \times n_w^{[l]} \times n_c^{[l]}$

$$n_H^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rfloor$$


$A^{[l]} : m \times n_H^{[l]} \times n_w^{[l]} \times n_c^{[l]}$

A Simple Convolution Network Example



A Simple Convolution Network Example

Example ConvNet



$$n_H^{[l]} = \left\lfloor \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right\rfloor$$

A simple convolution network example

Types of layer in a convolutional network

- Convolution (Conv net) Layer
- Pooling (Pool) Layer
- Fully connected (FC) Layer



Next:

Lab Practice Transfer Learning

