

# Mid-term Review and Final Presentation

11210IPT 553000

Deep Learning in Biomedical Optical Imaging

2023/11/20

# Mid-term Review



### Grading

- **Homework (35 %):** Assignments to get familiar with basic deep learning programming.
- **Midterm (30 %):** Basic concepts of deep learning and related mathematical derivation.
- **Final Project Presentation (20 %):** A project presentation of a research paper related to the deep learning application in medical imaging and the code implementation of the
- **Creativity Report (15 %):** A Provide a detailed analysis of your model implementation on a given image dataset.

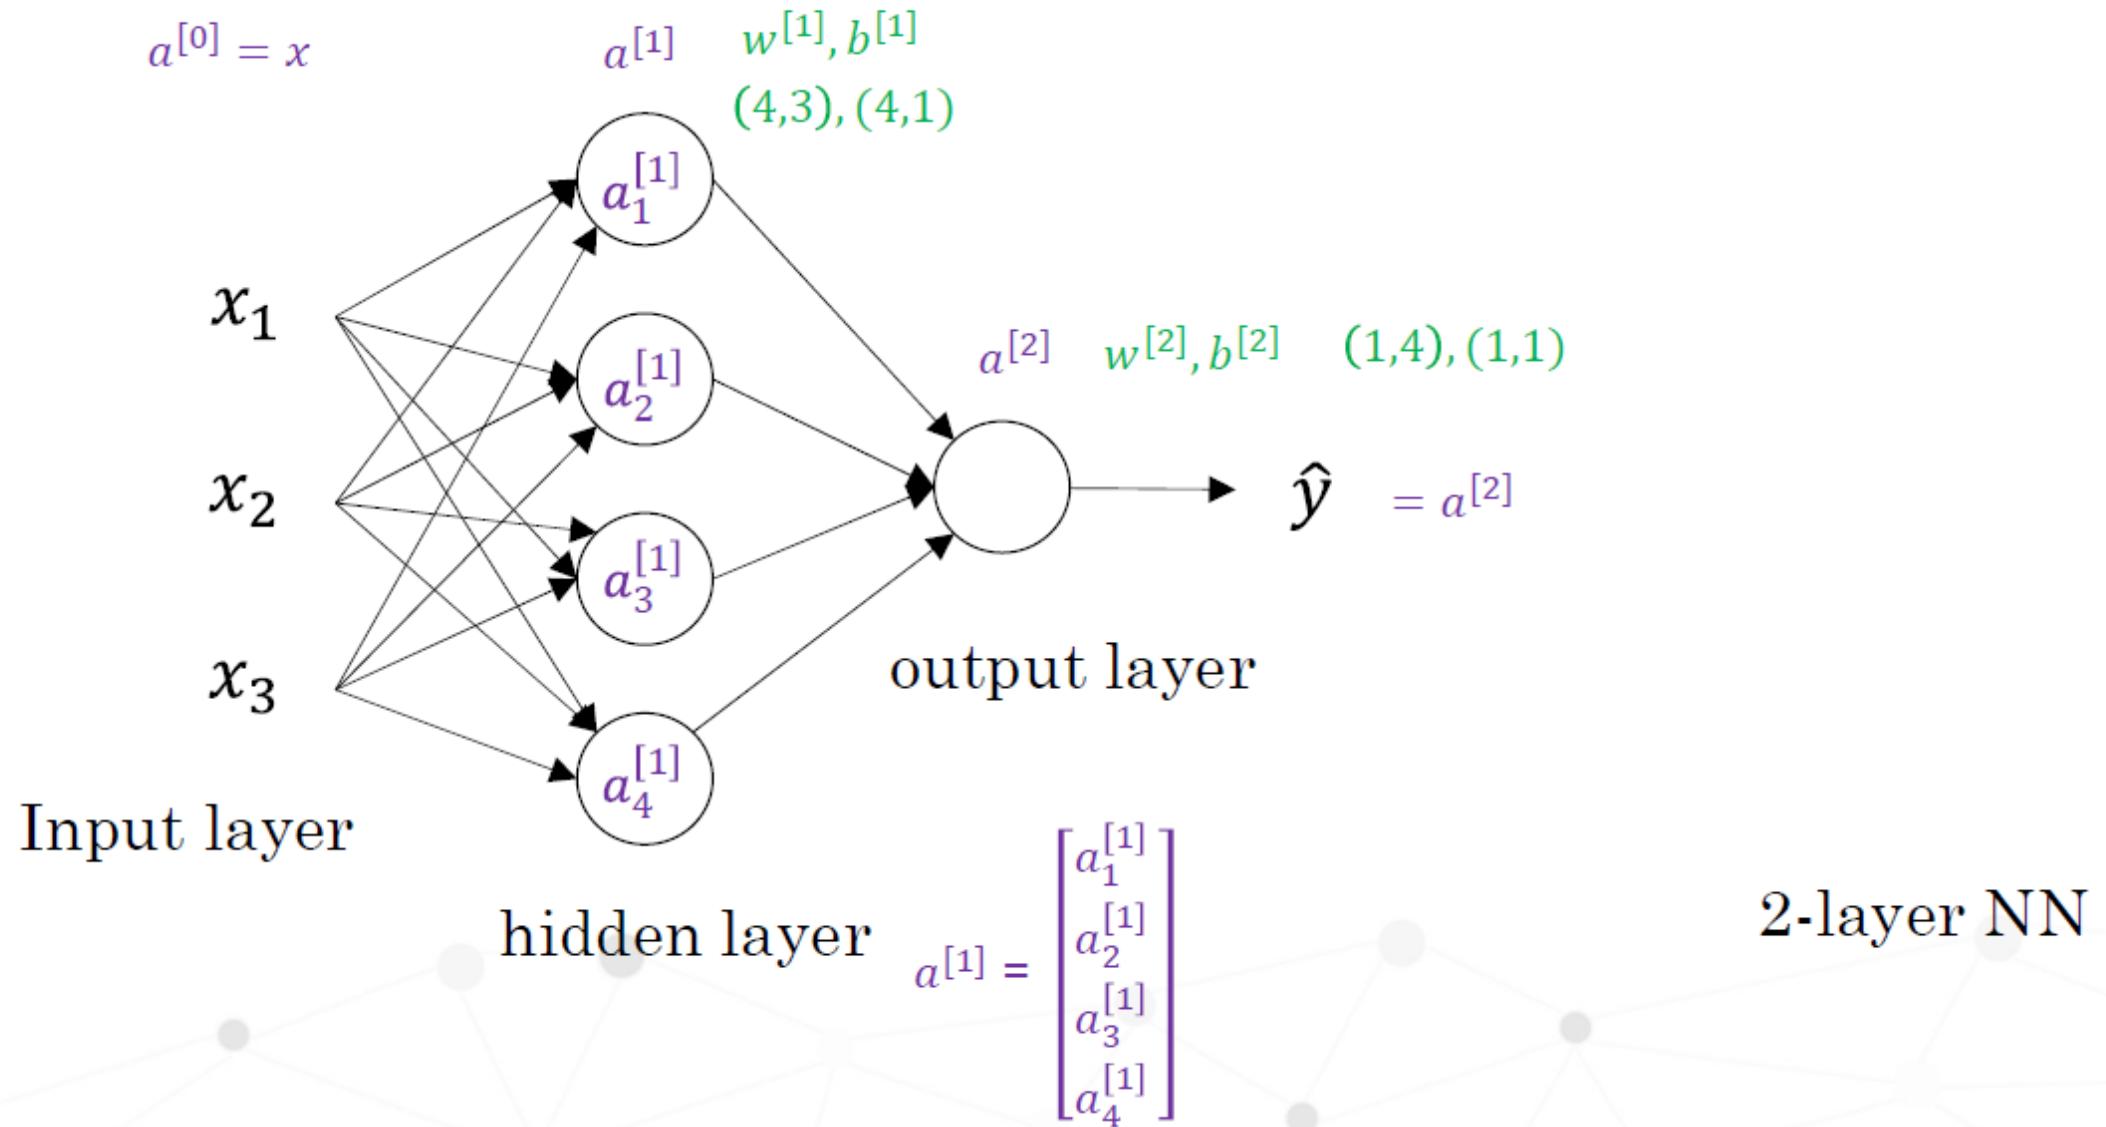
## Important Concepts

- Loss Function / Cost Function
- Activation functions
- Weight Initialization
- Forward propagation
- Backward propagation
- Hyperparameters
- Parameters
- Training, validation, and test datasets
- Bias and Variance
- Overfitting
- Regularization
- Dropout
- Data augmentation
- Early stopping
- Input normalization
- Vanishing/exploding gradients
- Epoch / Batch size/ Iteration
- Learning rate settings
- Saddle point
- Optimization algorithm: Momentum / RMSprop/ Adam
- Batch Normalization
- Multiclass classification
- Cross-entropy
- Transfer Learning
- Image Kernels (Filters)
- Convolution
- Pooling
- Padding
- Stride
- Features
- CNN networks
- Unsupervised Learning
- Recurrent-based model
- Attention-based model

- NumPy API
  - How to use some functions?
  - np.zeros, xxx.shape, np.add...

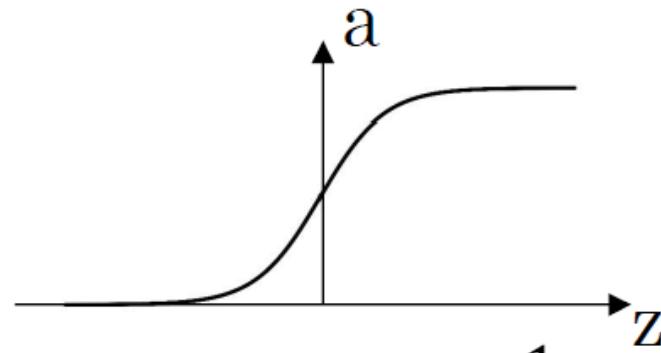
## One hidden layer Neural Network

### Neural Network Representation

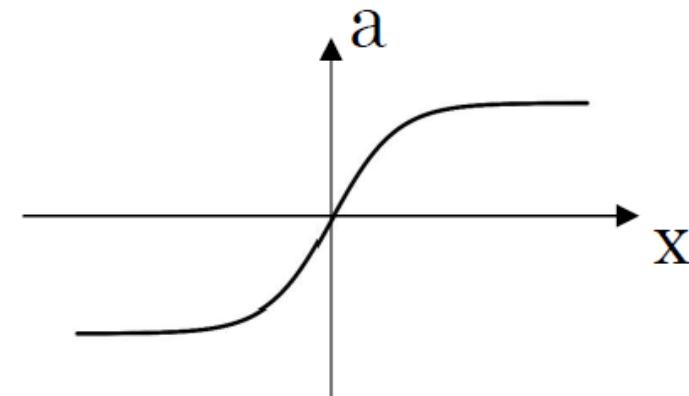


## Activation functions

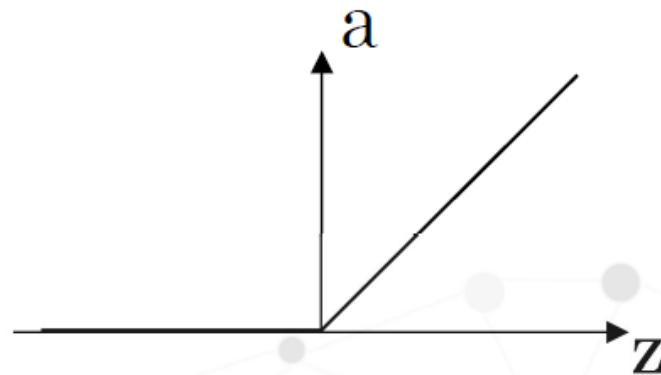
### Pros and cons of activation functions



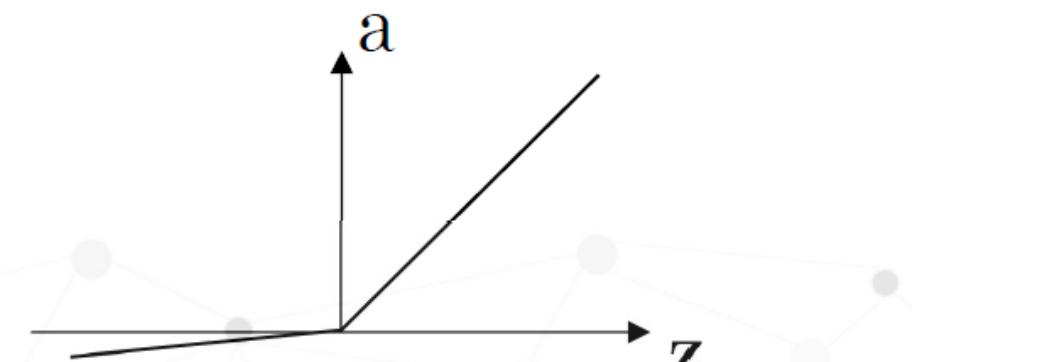
$$\text{sigmoid: } a = \frac{1}{1 + e^{-z}}$$



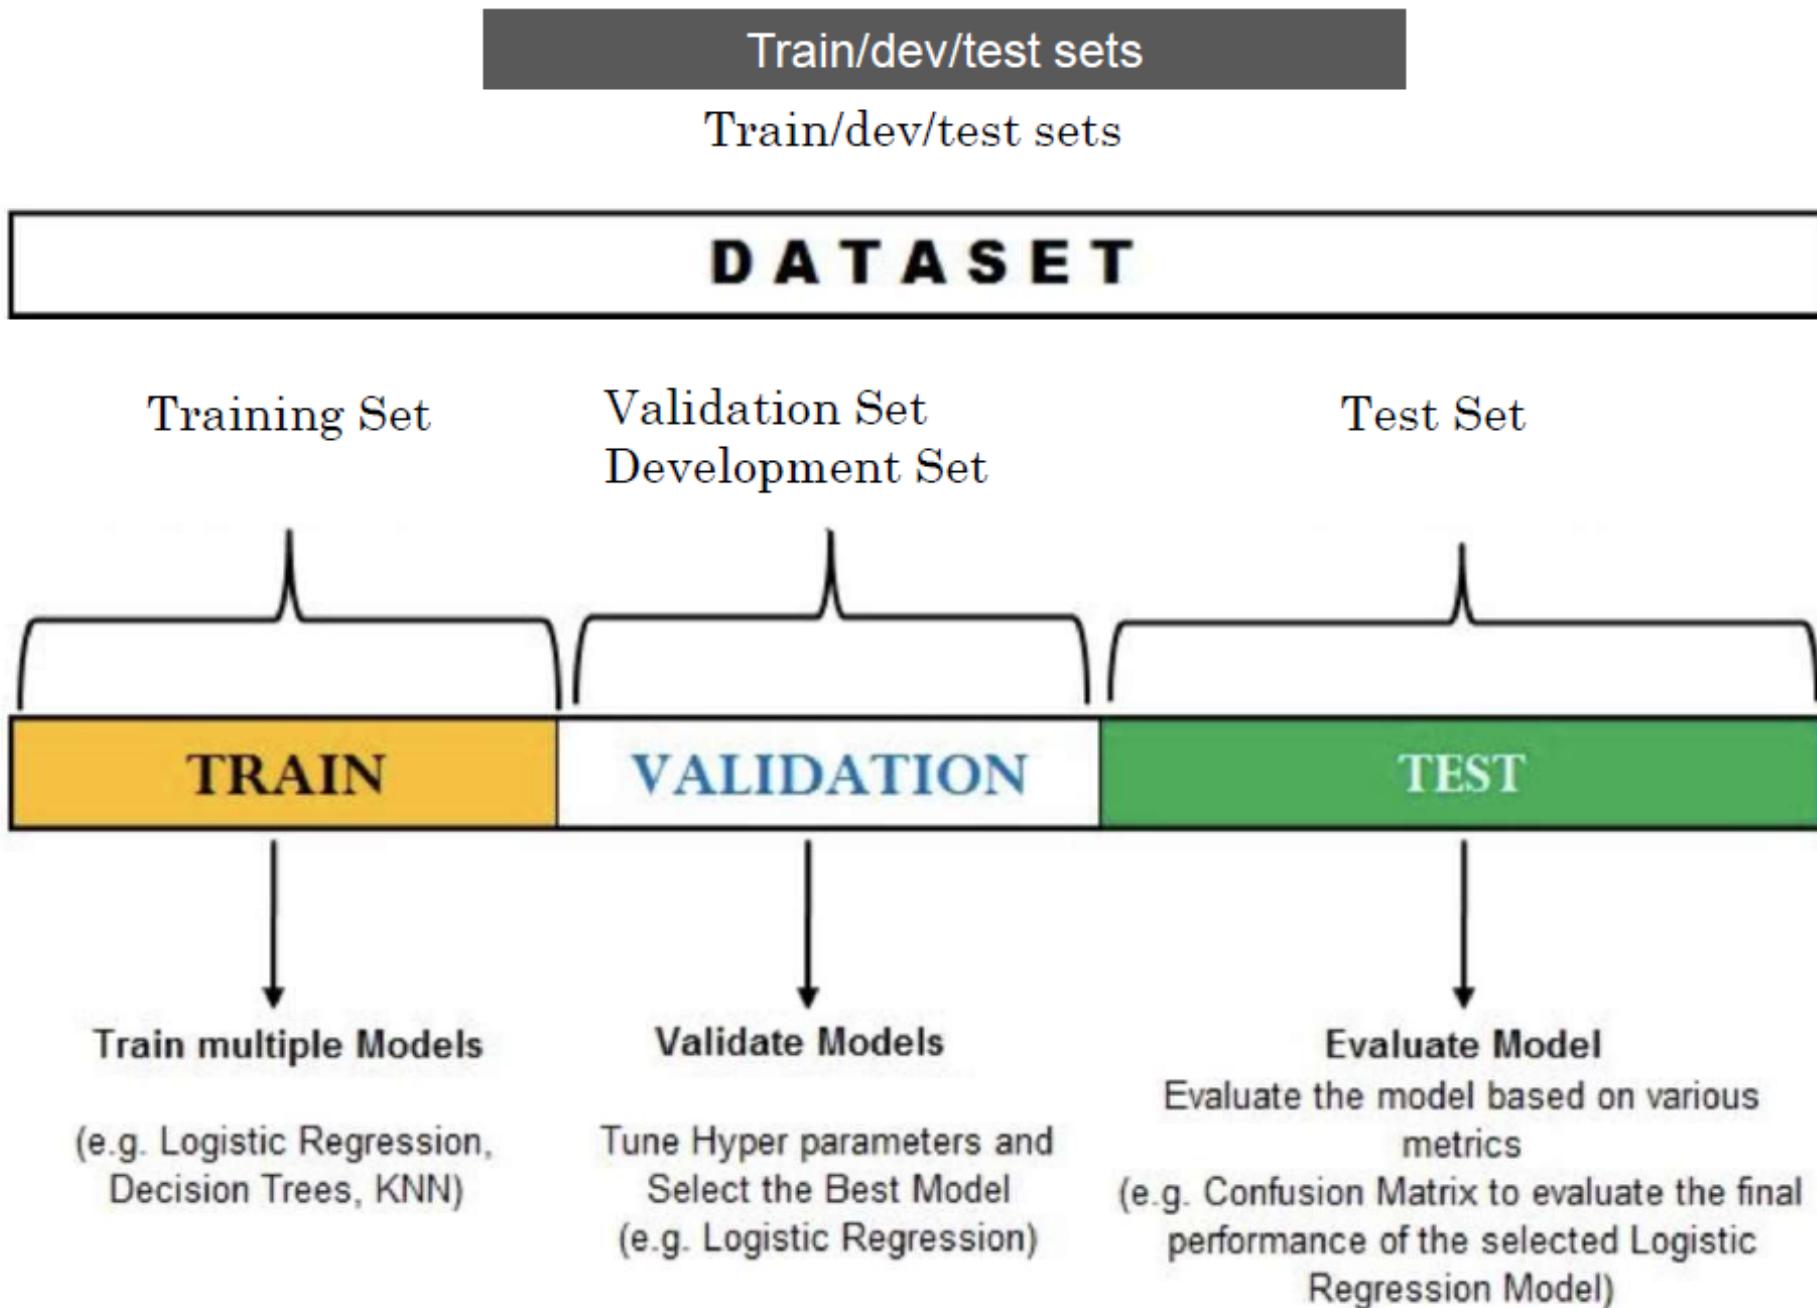
$$\tanh: a = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



$$\text{Relu: } a = \max(0, z)$$



$$\text{Leaky Relu: } a = \max(0.01z, z)$$



## Bias and Variance

## Cat classification

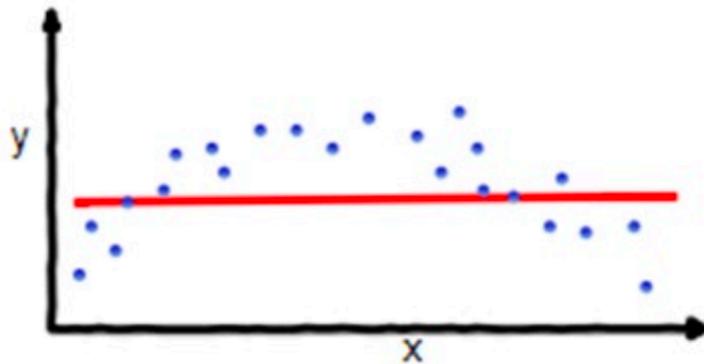


Train set error:	1%	15%	15%	0.5%
Dev set error:	11%	16%	30%	1%
	low bias high variance overfitting	high bias low variance underfitting	high bias high variance	low bias low variance

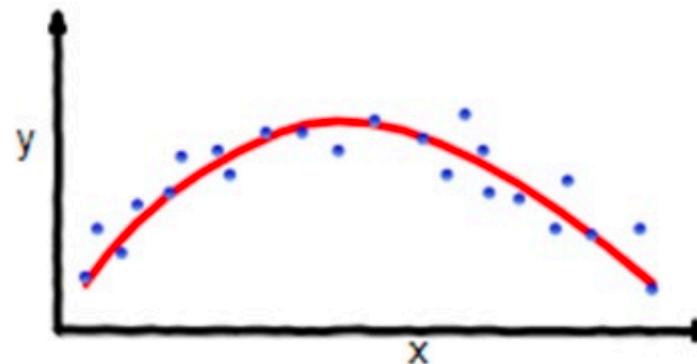
Human error (Bayes error): ~0%

### Bias and variance

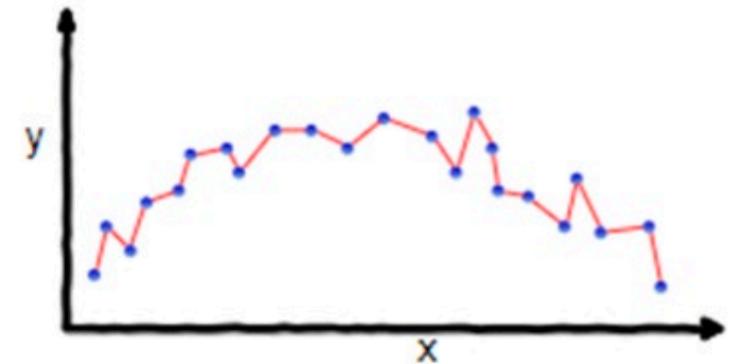
Model —————



high bias  
underfitting



“Just right”

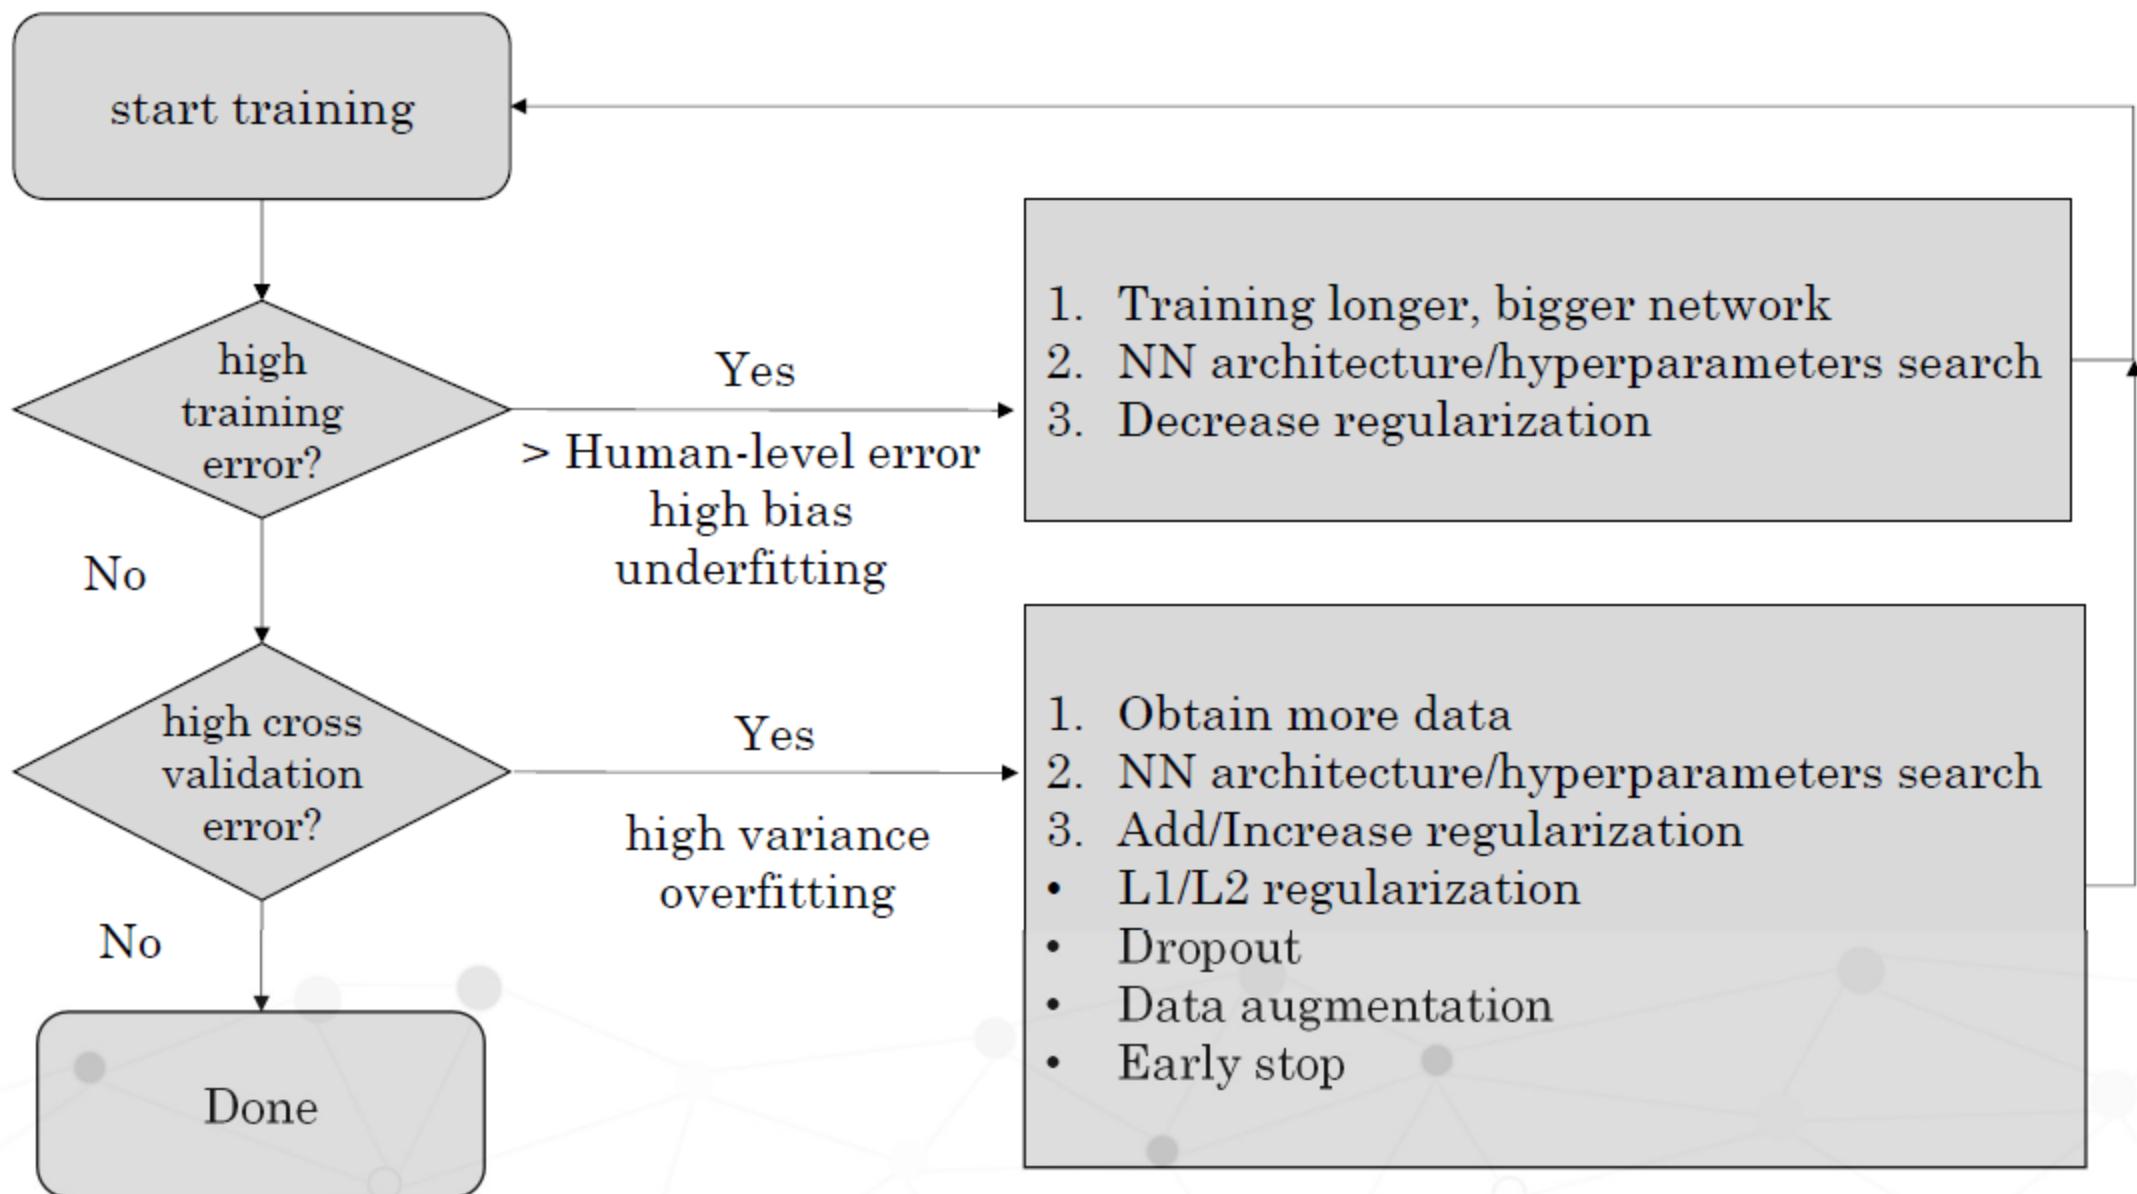


high variance  
overfitting



## Setting up your ML application

### Basic “recipe” for machine learning

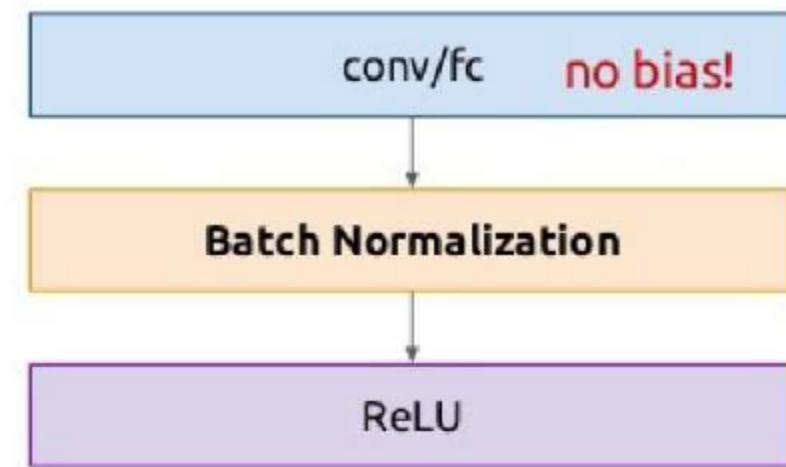


## Batch Normalization

Works by **re-normalizing layer inputs to have zero mean and unit standard deviation** with respect to running batch estimates.

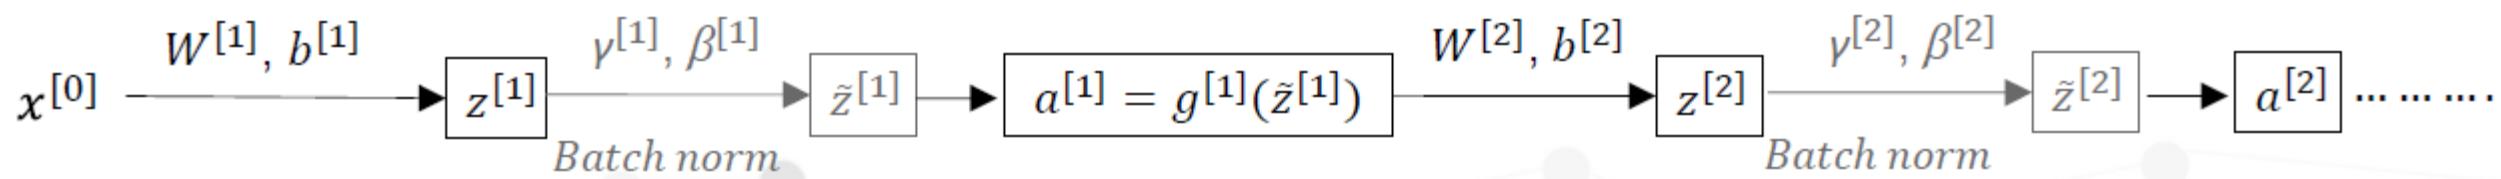
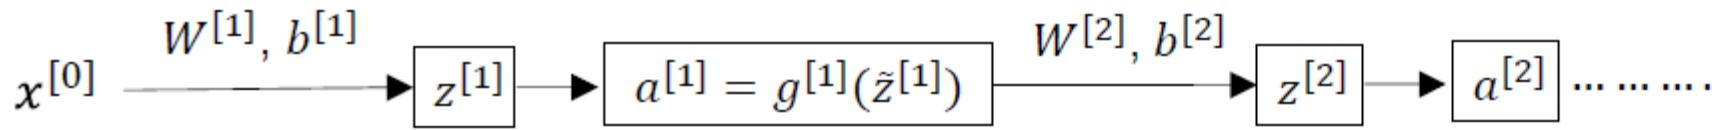
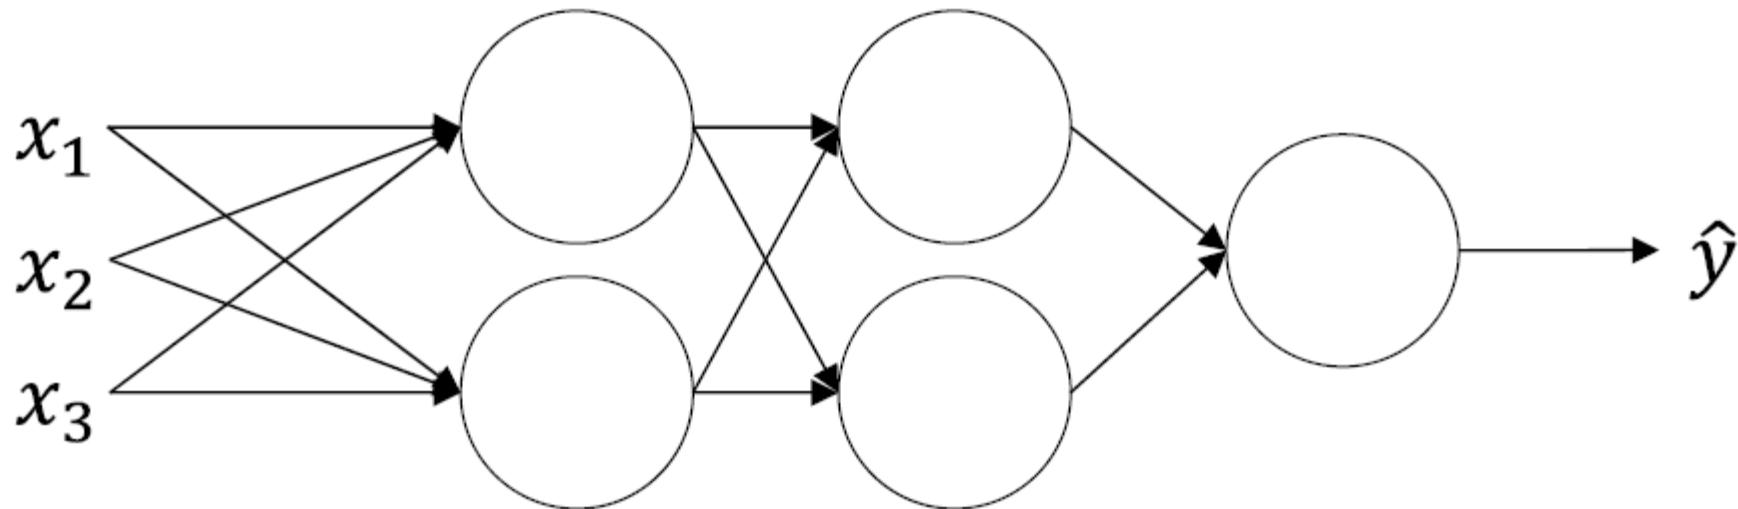
Also adds a **learnable scale and bias** term to allow the network to still use the nonlinearity.

Usually **allows much higher learning rates!**



## Fitting Batch Norm into a neural network

### Adding Batch Norm to a network



Parameters:  $W^{[1]}, b^{[1]}, W^{[2]}, b^{[2]}, \dots, W^{[L]}, b^{[L]}$   
 $\gamma^{[1]}, \beta^{[1]}, \gamma^{[2]}, \beta^{[2]}, \dots, \gamma^{[L]}, \beta^{[L]}$

## Vertical edge detection

$$3X1 + 1X1 + 1X2 + 0X0 + 0X5 + 0X7 + 1X-1 + 8X-1 + 2X-1$$

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

 $6 \times 6$ 

Filter or Kernel

1	0	-1
1	0	-1
1	0	-1

 $*$  $3 \times 3$ 

Convolution

-5	-4	0	8
-10	-2	2	3
0	-2	-4	-7
-3	-2	-3	-16

 $4 \times 4$

## Strided Convolutions

Output dimension

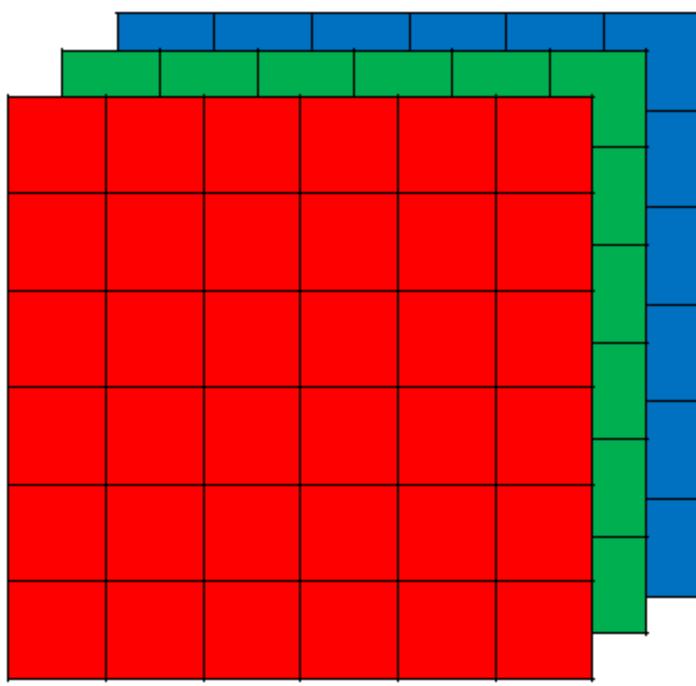
$n \times n$  image     $f \times f$  filter    padding  $p$     stride  $s$

Output size

$$\left[ \frac{n+2p-f}{s} + 1 \right] \quad \times \quad \left[ \frac{n+2p-f}{s} + 1 \right]$$

$$\lfloor x \rfloor = \text{floor}(x)$$

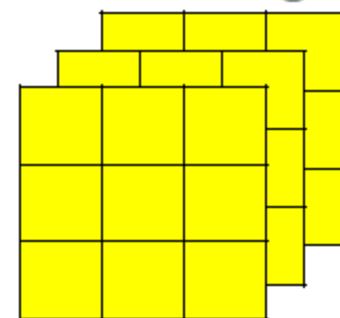
## Convolution Over Volumes



Multiple filters

Vertical edge

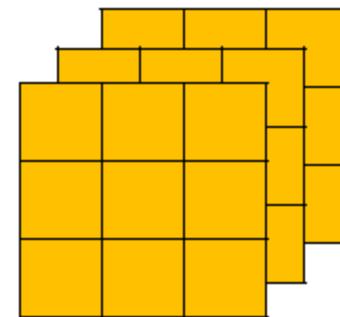
\*



$3 \times 3 \times 3$

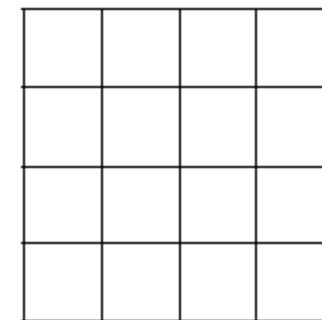
Horizontal edge

\*



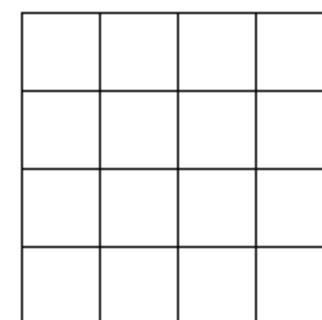
$3 \times 3 \times 3$

=



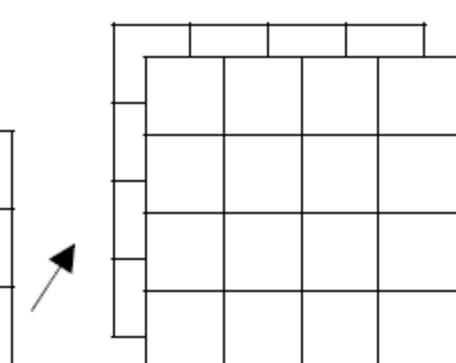
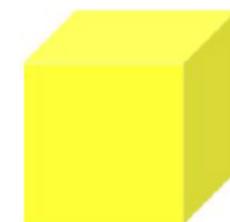
$4 \times 4$

=



$4 \times 4$

$4 \times 4 \times 2$



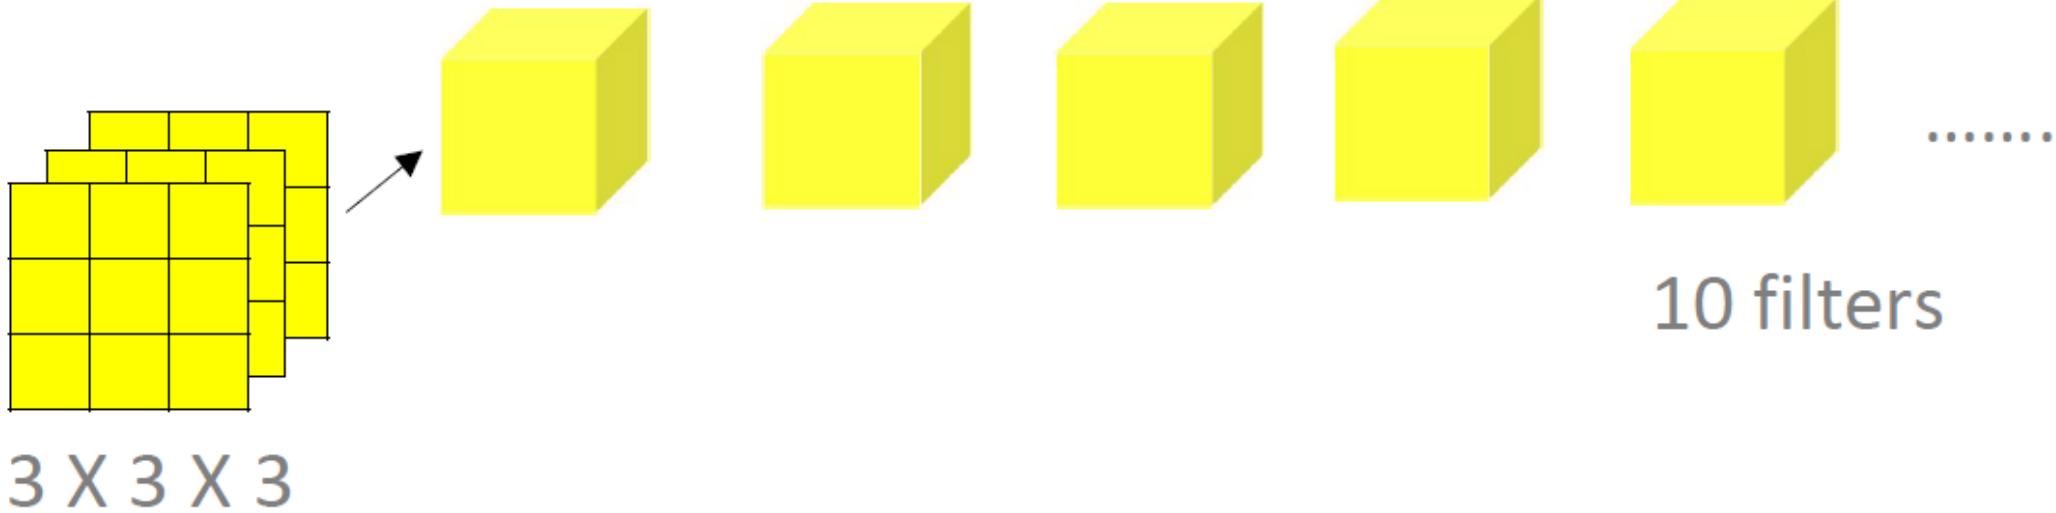
# of filters

Summary:  $n \times n \times n_c * f \times f \times n_c \rightarrow (n-f+1) \times (n-f+1) \times n'c$

## One Layer of a Convolutional Network

Number of parameters in one layer

If you have 10 filters that  $3 \times 3 \times 3$  in one layer of a neural network, how many parameters does this layer have?



$$\begin{aligned} & 27 \text{ parameters} + 1 \text{ bias} \\ & = 28 \text{ parameters} \end{aligned}$$

$$28 \times 10 = 280 \text{ parameters}$$

[Search Docs](#)[Community \[ + \]](#)[Developer Notes \[ + \]](#)[Language Bindings \[ + \]](#)[Python API \[ - \]](#)[torch](#)

### ● torch.nn

[torch.nn.functional](#)[torch.Tensor](#)[Tensor Attributes](#)[Tensor Views](#)[torch.amp](#)[torch.autograd](#)[torch.library](#)[torch.cpu](#)[torch.cuda](#)[Understanding CUDA Memory Usage](#)[Generating a Snapshot](#)[Using the visualizer](#)

## CONV2D

[Conv2d](#)[Conv2d](#)

CLASS `torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1, groups=1, bias=True, padding_mode='zeros', device=None, dtype=None)` [\[SOURCE\]](#) <https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C_{\text{in}}, H, W)$  and output  $(N, C_{\text{out}}, H_{\text{out}}, W_{\text{out}})$  can be precisely described as:

$$\text{out}(N_i, C_{\text{out}_j}) = \text{bias}(C_{\text{out}_j}) + \sum_{k=0}^{C_{\text{in}}-1} \text{weight}(C_{\text{out}_j}, k) \star \text{input}(N_i, k)$$

where  $\star$  is the valid 2D [cross-correlation](#) operator,  $N$  is a batch size,  $C$  denotes a number of channels,  $H$  is a height of input planes in pixels, and  $W$  is width in pixels.

This module supports [TensorFloat32](#).

On certain ROCm devices, when using float16 inputs this module will use [different precision](#) for backward.

- `stride` controls the stride for the cross-correlation, a single number or a tuple.
- `padding` controls the amount of padding applied to the input. It can be either a string {'valid', 'same'} or an int / a tuple of ints giving the amount of implicit padding applied on both sides.
- `dilation` controls the spacing between the kernel points; also known as the  $\texttt{à trous}$  algorithm. It is

[Search Docs](#)[Community \[ + \]](#)[Developer Notes \[ + \]](#)[Language Bindings \[ + \]](#)[Python API \[ - \]](#)[torch](#)[● torch.nn](#)[torch.nn.functional](#)[torch.Tensor](#)[Tensor Attributes](#)[Tensor Views](#)[torch.amp](#)[torch.autograd](#)[torch.library](#)[torch.cpu](#)[torch.cuda](#)[Understanding CUDA Memory Usage](#)[Generating a Snapshot](#)

## MAXPOOL2D

CLASS `torch.nn.MaxPool2d(kernel_size, stride=None, padding=0, dilation=1, return_indices=False, ceil_mode=False)` [\[SOURCE\]](#)



Applies a 2D max pooling over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size  $(N, C, H, W)$ , output  $(N, C, H_{out}, W_{out})$  and `kernel_size`  $(kH, kW)$  can be precisely described as:

$$\begin{aligned} \text{out}(N_i, C_j, h, w) = & \max_{m=0, \dots, kH-1} \max_{n=0, \dots, kW-1} \\ & \text{input}(N_i, C_j, \text{stride}[0] \times h + m, \text{stride}[1] \times w + n) \end{aligned}$$

If `padding` is non-zero, then the input is implicitly padded with negative infinity on both sides for `padding` number of points. `dilation` controls the spacing between the kernel points. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.

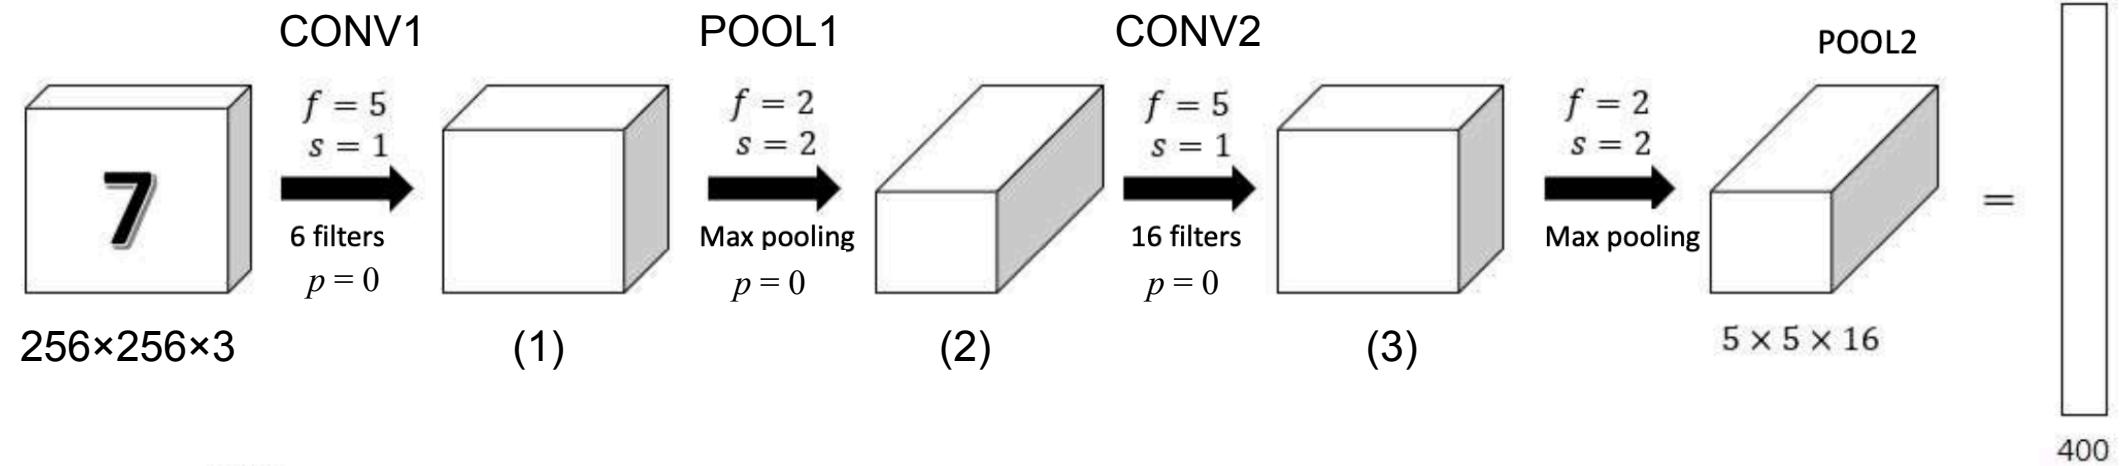
<https://pytorch.org/docs/stable/generated/torch.nn.MaxPool2d.html>

• NOTE

When `ceil_mode=True`, sliding windows are allowed to go off-bounds if they start within the left padding or the input. Sliding windows that would start in the right padded region are ignored.

## Convolutional neural network example

### Neural network example (~LeNet-5)



## Key evaluation metrics

### Confusion matrix

#### Model Output

		+	-
GT	Disease	True Positive (TP)	False Negative (FN)
	Normal	False Positive (FP)	True Negative (TN)

$$\xrightarrow{\hspace{1cm}} \frac{\#(+ \text{ and disease})}{\#(\text{disease})} = \text{Sensitivity}$$

$$\xrightarrow{\hspace{1cm}} \frac{\#(- \text{ and normal})}{\#(\text{normal})} = \text{Specificity}$$

$$\downarrow \qquad \qquad \qquad \text{PPV} = \frac{\#(+ \text{ and disease})}{\#(+)}$$

$$\downarrow \qquad \qquad \qquad \text{NPV} = \frac{\#(- \text{ and normal})}{\#(-)}$$

# Final Presentation

- Midterm Review on 11/20 and Midterm on 11/27
- Research Presentation Topic
  - Presentation Week (12/18, 12/25, 1/8)
  - Volunteers of the first week and randomly pick the presentation order on 11/20
  - Pick a research paper published after 2021
  - Send the title and the paper file to TA ASAP
  - Get the confirmation from the lecturer **due on 12/4**

- **Abstract**  
What is this research about?
- **Motivation / Purpose**  
Why do the authors perform this research?
- **Introduction to the Bio-medical Imaging Technology**
- **Network Architecture**  
Introduce the deep learning model including its design, pros and cons
- **Comparison**  
Compare the performance of this research to other researches or baselines
- **Conclusion**  
Your conclusion
- **Code Implementation (Bonus)**  
Show your results after running the codes (Submit the files or the link to TA)  
Modifications are strongly encouraged!

When seeking research papers for your presentation, consider those published in renowned journals or international conferences, such as Nature, IEEE journals, OPTICA, or conferences like CVPR, ICCV, ECCV, NeurIPS, ICML, ICLR.

[https://scholar.google.com/citations?view\\_op=top\\_venues&hl=zh-TW](https://scholar.google.com/citations?view_op=top_venues&hl=zh-TW)

文章	H5 指數	H5 中位數
1. Nature	467	707
2. The New England Journal of Medicine	439	876
3. Science	424	665
4. IEEE/CVF Conference on Computer Vision and Pattern Recognition	422	681
5. The Lancet	368	688
6. Nature Communications	349	456
7. Advanced Materials	326	415
8. Cell	316	503
9. Neural Information Processing Systems	309	503
10. International Conference on Learning Representations	303	563
11. JAMA	286	476
12. Science of The Total Environment	273	375
13. Nature Medicine	268	459
14. Proceedings of the National Academy of Sciences	268	394
15. Angewandte Chemie International Edition	266	362
16. Chemical Reviews	264	459

類別 > Engineering & Computer Science > Computer Vision & Pattern Recognition ▾

文章	H5 指數	H5 中位數
1. IEEE/CVF Conference on Computer Vision and Pattern Recognition	422	681
2. European Conference on Computer Vision	238	390
3. IEEE/CVF International Conference on Computer Vision	228	366
4. IEEE Transactions on Pattern Analysis and Machine Intelligence	179	318
5. IEEE Transactions on Image Processing	138	199
6. Pattern Recognition	111	160
7. IEEE/CVF Computer Society Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)	108	176
8. Medical Image Analysis	103	171
9. IEEE/CVF Winter Conference on Applications of Computer Vision (WACV)	95	150
10. International Journal of Computer Vision	88	165
11. Pattern Recognition Letters	80	138
12. British Machine Vision Conference (BMVC)	77	128
13. IEEE/CVF International Conference on Computer Vision Workshops (ICCVW)	66	102
14. IEEE International Conference on Image Processing (ICIP)	61	90
15. Asian Conference on Computer Vision (ACCV)	60	92

# Example

Published as a conference paper at ICLR 2021

## FULLY UNSUPERVISED DIVERSITY DENOISING WITH CONVOLUTIONAL VARIATIONAL AUTOENCODERS

**Mangal Prakash \***

Center for Systems Biology Dresden  
Max-Planck Institute (CBG)  
Dresden, Germany  
prakash@mpi-cbg.de

**Alexander Krull \*†**

School of Computer Science  
University of Birmingham  
Birmingham, UK  
a.f.f.krull@bham.ac.uk

**Florian Jug<sup>†</sup>**

Center for Systems Biology Dresden  
Max-Planck Institute (CBG)  
Dresden, Germany  
Fondazione Human Technopole, Milano, Italy  
jug@mpi-cbg.de, florian.jug@fht.org

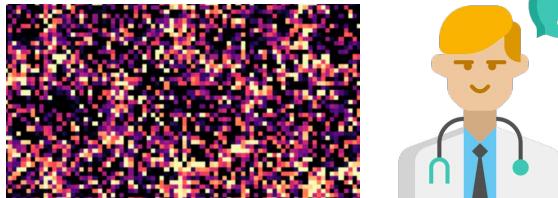
### ABSTRACT

Deep Learning based methods have emerged as the indisputable leaders for virtually all image restoration tasks. Especially in the domain of microscopy images, various content-aware image restoration (CARE) approaches are now used to improve the interpretability of acquired data. Naturally, there are limitations to what can be restored in corrupted images, and like for all inverse problems, many potential solutions exist, and one of them must be chosen. Here, we propose DIVNOISING, a denoising approach based on fully convolutional variational autoencoders (VAEs), overcoming the problem of having to choose a single solution by predicting a whole distribution of denoised images. First we introduce a principled way of formulating the unsupervised denoising problem within the VAE framework by explicitly incorporating imaging noise models into the decoder. Our approach is fully unsupervised, only requiring noisy images and a suitable description of the imaging noise distribution. We show that such a noise model can either be measured, bootstrapped from noisy data, or co-learned during training. If desired, consensus predictions can be inferred from a set of DIVNOISING predictions, leading to competitive results with other unsupervised methods and, on occasion, even with the supervised state-of-the-art. DIVNOISING samples from the posterior enable a plethora of useful applications. We are (i) showing denoising results for

# Introduction

## Image Restoration

- The goal of scientific image analysis is to analyze pixel-data and measure the properties of objects of interest in images.
- Image restoration is the task of removing unwanted noise and distortions, giving us clean images that are closer to the true but unknown signal.

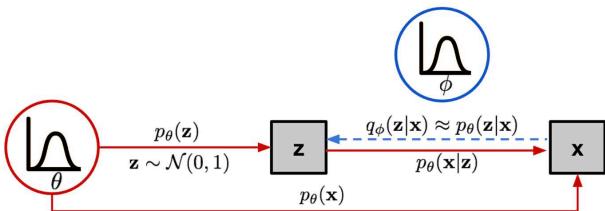
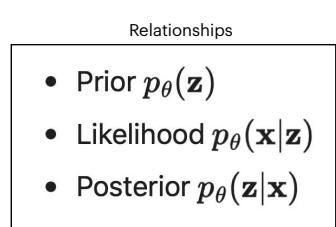


1. Begin by providing an overview of the background and existing research in the paper's domain.

# Introduction

## Variational Autoencoder (VAE)

- VAEs are generative models, capable of learning complex distributions over images.



Instead of mapping the input into a *fixed* vector, VAE wants to map it into a distribution.

Fig. 6. The graphical model involved in Variational Autoencoder. Solid lines denote the generative distribution  $p_\theta(\cdot)$  and dashed lines denote the distribution  $q_\phi(\cdot|\mathbf{x})$  to approximate the intractable posterior  $p_\theta(\cdot|\mathbf{x})$ .

[source](#)

# Introduction

## Motivations

- Distortions degrade some of the information content in images, generally making it impossible to fully recover the desired clean signal with certainty.
- Generative models, such as VAEs, have been overlooked as a method to solve unsupervised image denoising problems.
  - This might also be due to the fact that vanilla VAEs show sub-par performance on denoising problems.
- VAEs are a canonical choice when a distribution over a set of variables needs to be learned.

2. Explain the challenges or problems the authors aim to address. Highlight the key contributions of the paper.

## Introduction

### Motivations

- Distortions degrade some of the information content in images, generally making it impossible to fully recover the desired clean signal with certainty.
- Generative models, such as VAEs, have been overlooked as a method to solve unsupervised image denoising problems.
  - This might also be due to the fact that vanilla VAEs show sub-par performance on denoising problems.
  - VAEs are a canonical choice when a distribution over a set of variables needs to be learned.

12

## Introduction

### Contributions

- DivNoisig will be hugely beneficial for computational biology applications in biomedical imaging, where noise is typically unavoidable and huge datasets need to be processed on a daily basis.
- Enable unsupervised diverse SOTA denoising while requiring only comparatively little computational resources, rendering our approach particularly practical.
- Have the potential to be useful for many real-world applications and will not only generate state-of-the-art (SOTA) restored images, but also enrich quantitative downstream processing.

15

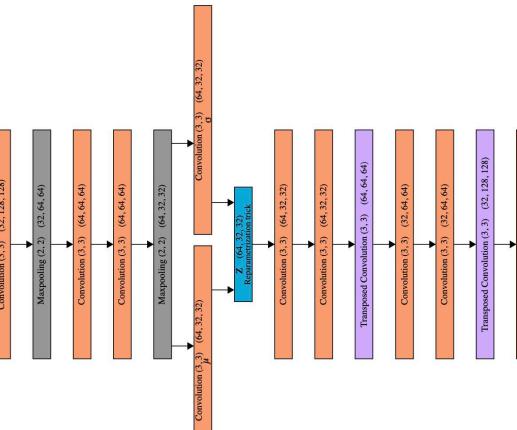
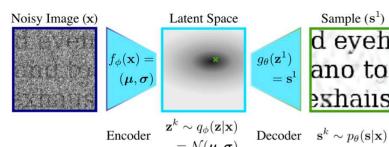
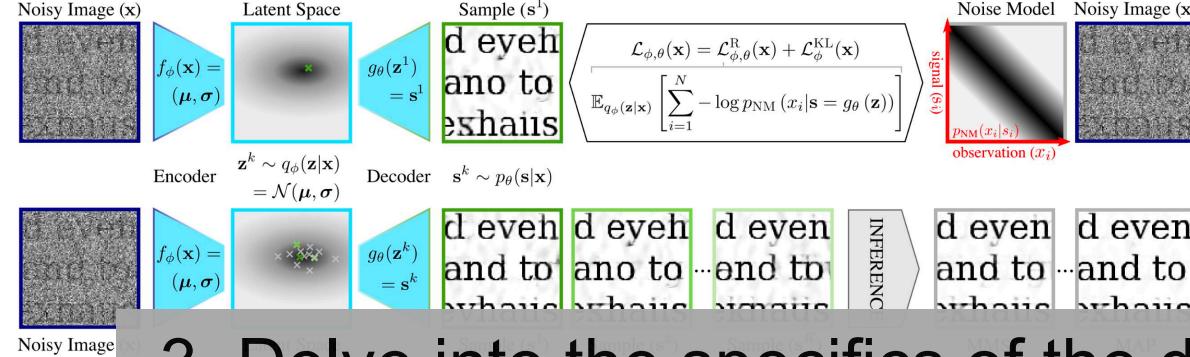
28

# Methods

## Network Architecture

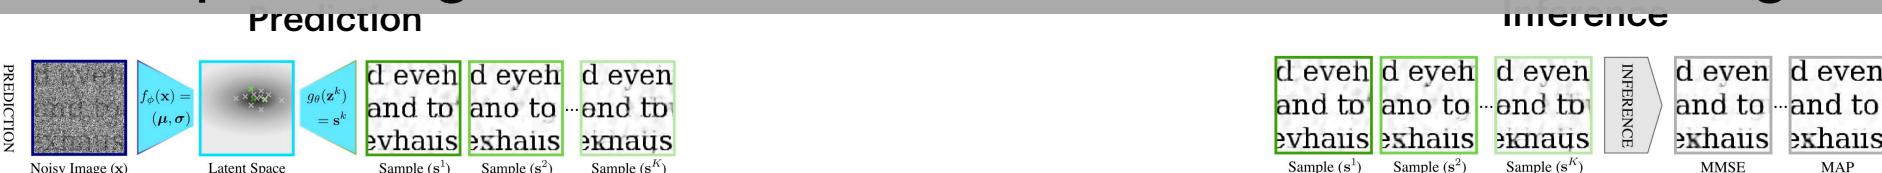
# Methods

## DivNoising



These networks count about 200k parameters and have a GPU memory footprint of approximately 1.6GB on an NVIDIA TITAN Xp.

3. Delve into the specifics of the deep learning model, including hyper-parameters and loss functions. Discuss the fundamental theories or unique designs the authors used to tackle their challenges.



- Get samples from an approximate posterior
  - Feed the noisy image into our encoder.
  - Draw samples  $\mathbf{z}^k \sim q_\phi(\mathbf{z}|\mathbf{x})$
  - Decode the samples via the decoder to get  $\mathbf{s}^k = g_\theta(\mathbf{z}^k)$

- Given a set of posterior samples for a noisy image, we can infer different point estimates.
  - Approximate the MMSE estimate by averaging many samples.
  - Find the *maximum a posteriori* (MAP) estimate.
    - i.e. the most likely signal given the noisy observation, by finding the mode of the posterior distribution. For this purpose, we iteratively use the mean shift algorithm with decreasing bandwidth to find the mode of our sample set.

4. Describe the dataset used, including the number of images, the imaging technology, and whether the dataset is private or public.

## Results

### Dataset

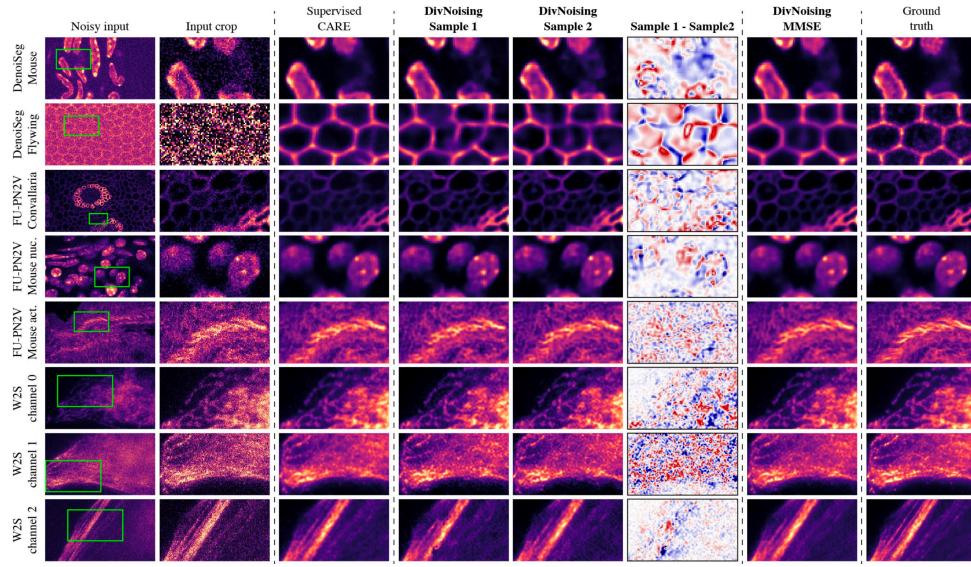
- **Intrinsically Noisy data:** No prior noise model is known for microscopy datasets.
- For these datasets, we used GMM-based noise models, which are measured from calibration images, as well as co-learned noise models.
  - The *Convallaria* and mouse actin datasets are acquired on a **spinning disc confocal microscope** while the mouse skull nuclei dataset is acquired with a **point scanning confocal microscope**.
  - Widefield2SIM (W2S), acquired using a **conventional fluorescence widefield and SIM imaging**. W2S includes 144,000 real fluorescence microscopy images, resulting in a total of 360 sets of images.

## Results

### Dataset

- **Synthesized noise dataset**
  - *DenoiSeg Mouse* data, showing cell nuclei in the developing mouse skull, consists of 908 training and 160 validation images of size 128x128, with additional 67 images of size 256x256 for testing.
    - Two noisy datasets were created with this data.
    - Exposing all images to pixel-wise independent Gaussian noise
    - First, apply poisson noise and followed by adding gaussian noise (randomly change 3% of pixels to either 0 or 255)
  - *DenoiSeg Flywing* data is showing membrane labeled cells in a fly wing, consisting of 1428 training and 252 validation patches of size 128x128, with additional 42 images of size 512x512 for testing.

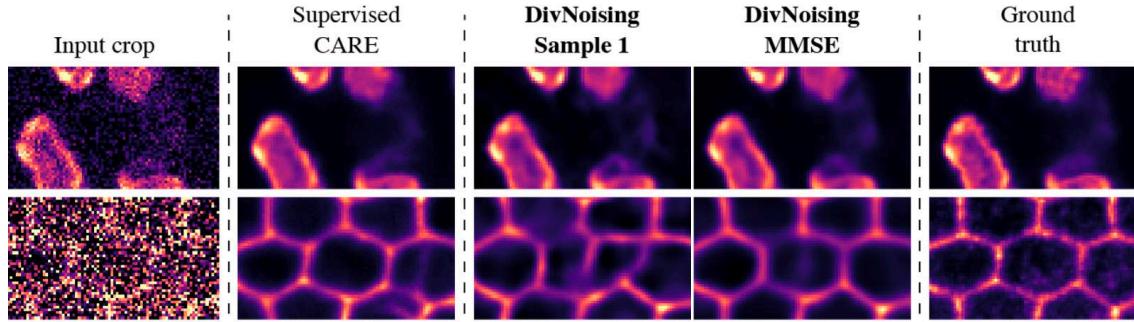
# Results



38

5. Present both quantitative and qualitative results. Clarify the evaluation metrics used and explain the significance of the images and data presented.

## Results Qualitative - Synthesized Noise Dataset



39

## Results Quantitative

Dataset	Fully Unsupervised			Unsup. ( $p_{NM}$ requ.)		Supervised CARE
	N2V	Vanilla VAE	DivNoising	PN2V	DivNoising	
Convallaria	$35.73 \pm 0.037$	$36.57 \pm 0.033$	$36.78 \pm 0.007$	$36.47 \pm 0.031$	$36.70 \pm 0.012$	$36.90 \pm 0.004$
↓ Bootstrapped						$36.64 \pm 0.023$
Mouse Act.	$33.39 \pm 0.014$	$33.46 \pm 0.158$	$33.82 \pm 0.006$	$33.86 \pm 0.018$	$33.99 \pm 0.004$	$34.20 \pm 0.021$
Mouse Nuc.	$35.84 \pm 0.015$	$35.84 \pm 0.023$	$36.05 \pm 0.052$	$36.35 \pm 0.018$	$36.26 \pm 0.047$	$36.58 \pm 0.019$
Ch.0 (avg1)	$34.59 \pm 0.041$	$33.02 \pm 0.147$	$34.24 \pm 0.006$	-	$34.13 \pm 0.002$	$35.22 \pm 0.069$
Ch.1 (avg1)	$32.11 \pm 0.030$	$31.36 \pm 0.041$	$32.22 \pm 0.021$	-	$32.22 \pm 0.013$	$32.88 \pm 0.021$
Ch.2 (avg1)	$35.04 \pm 0.073$	$33.72 \pm 0.187$	$35.24 \pm 0.028$	$32.79 \pm 0.085$	$35.18 \pm 0.020$	$35.91 \pm 0.030$
W2S						
Ch.0 (avg16)	$39.01 \pm 0.019$	$39.27 \pm 0.192$	$39.45 \pm 0.036$	$39.36 \pm 0.103$	$39.63 \pm 0.007$	$42.35 \pm 0.012$
Ch.1 (avg16)	$37.91 \pm 0.059$	$38.33 \pm 0.021$	$38.41 \pm 0.018$	$38.46 \pm 0.012$	$38.39 \pm 0.007$	$39.64 \pm 0.061$
Ch.2 (avg16)	$40.30 \pm 0.023$	$40.24 \pm 0.043$	$40.56 \pm 0.019$	$40.36 \pm 0.091$	$40.41 \pm 0.041$	$42.03 \pm 0.027$
↓ Denoiseg						
Mouse	$33.84 \pm 0.070$	$34.06 \pm 0.003$	$34.06 \pm 0.005$	$34.19 \pm 0.037$	$34.13 \pm 0.003$	$35.11 \pm 0.016$
Flywing	$24.79 \pm 0.034$	$24.88 \pm 0.045$	$24.92 \pm 0.016$	$24.85 \pm 0.036$	$25.02 \pm 0.024$	$25.79 \pm 0.014$
Mouse s&p	$32.98 \pm 0.020$	$23.62 \pm 0.084$	$35.19 \pm 0.030$	$29.67 \pm 0.079$	$36.21 \pm 0.015$	$37.03 \pm 0.016$
BioID Face	$32.34 \pm 0.080$	$32.58 \pm 0.022$	$33.02 \pm 0.020$	$33.76 \pm 0.079$	$33.12 \pm 0.039$	$35.06 \pm 0.051$

43

6. Offer your perspective on the paper. Assess the feasibility and effectiveness of the proposed methods.

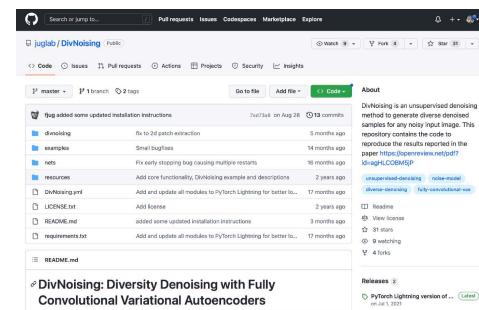
!!!!



**Code Implementation (optional):** If the authors have provided open-source code, consider trying it out or incorporating their ideas into your creative report!

## Code Implementation Overview

- Github: <https://github.com/juglab/DivNoising>
- I follow the official instructions to run ipython notebooks.
- 1. Training
- 2. Testing



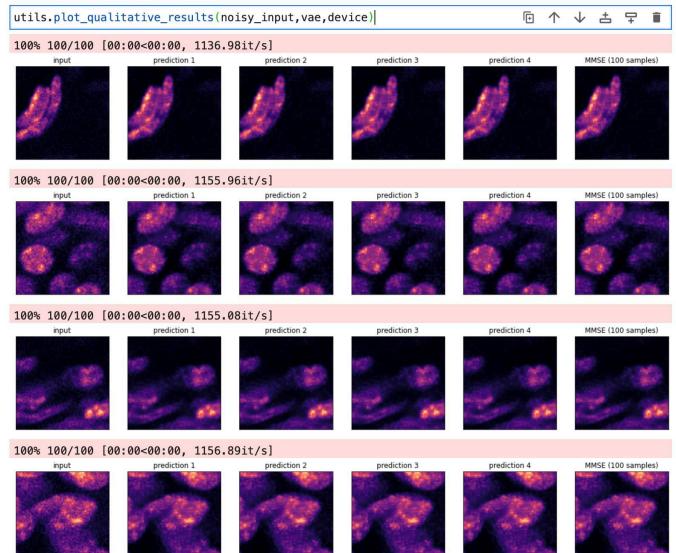
46

## Code Implementation Testing

- Load data to predict.
- Use trained model to generate samples.
  - MMSE samples: 10 (original setting is 1000.)
- And compute PSNR.

Dataset	Fully Unsupervised			Unsup. ( $p_{NM}$ requ.)	
	N2V	Vanilla VAE	DivNoising	PN2V	DivNoising
Convallaria	$35.73 \pm 0.037$	$36.57 \pm 0.033$	$36.78 \pm 0.007$	$36.47 \pm 0.031$	$36.90 \pm 0.004$
Bootstrapped				$36.70 \pm 0.012$	$36.64 \pm 0.023$

```
image: 198 psnr:32.877 mean psnr:32.818
image: 199 psnr:32.829 mean psnr:32.818
mean 32.81795153617859
```



48

33